


Real-time Neural Network Inference on Extremely Weak Devices: Agile Offloading with Explainable AI


Kai Huang, Wei Gao
University of Pittsburgh, USA

Real-time NN demands on weak embedded devices



Resources limitation on weak embedded devices



Weak Embedded Devices



<1MB memory and storage
16~216 MHz CPU

Large Neural Networks

Residual Networks (ResNet50)



Require **>100** MB of memory space
>2 GHz CPU for **60 ms** latency

Existing Work

- **Local Inference**
 - Prune, Compress, Network Architecture Search
 - Lead to **oversimplified NN structures**
 - **>10%** accuracy lost
- **Remote Inference**
 - Compress raw data before transmission
 - Limited data compressibility when accuracy loss is minimum.
- **NN offloading**
 - Use a local NN to sparsity & compress data
 - Higher compressibility but expensive local NNs
 - Extra transmission cost

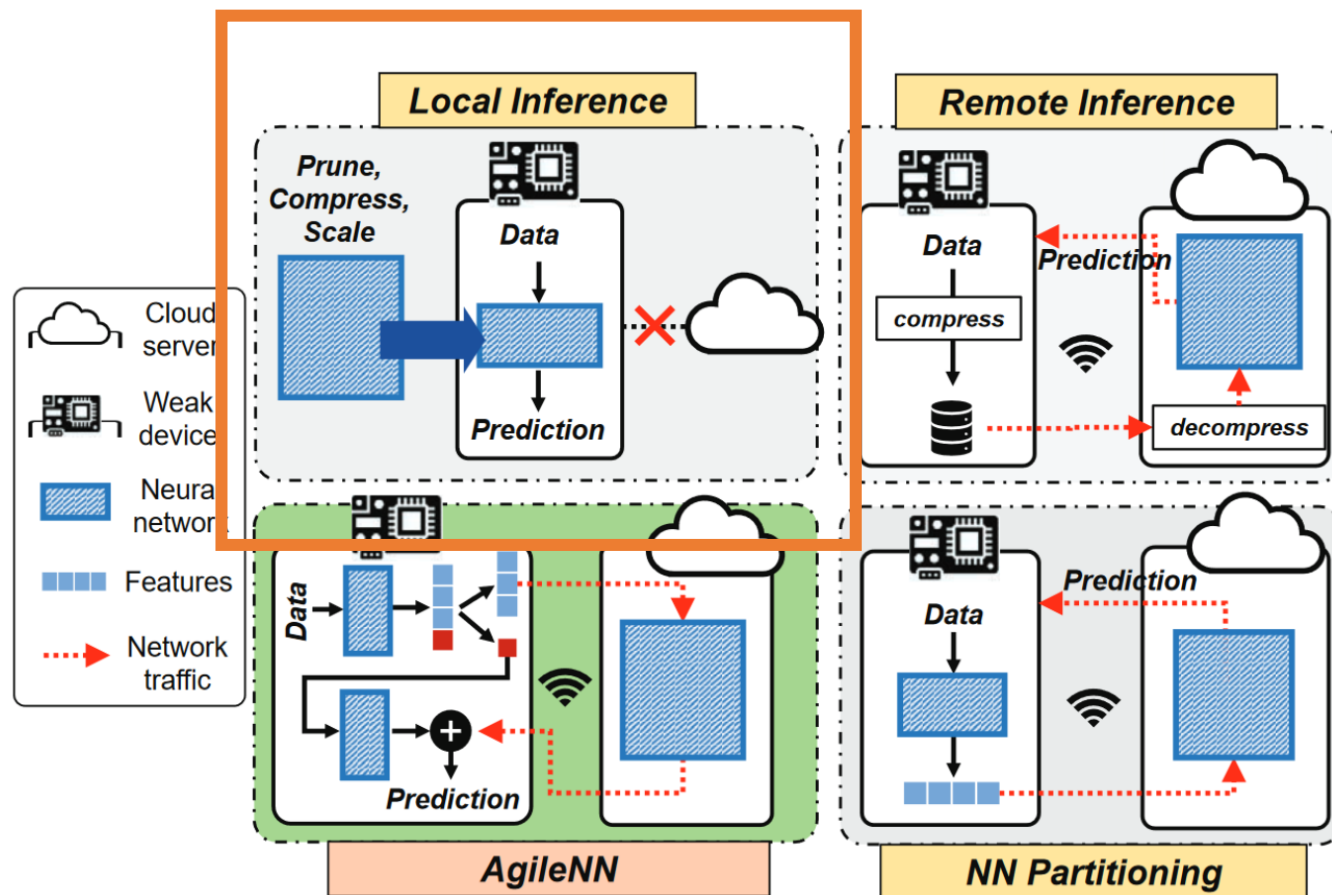


Figure 1: Existing work vs. AgileNN

Existing Work

- Local Inference
 - Prune, Compress, Network Architecture Search
 - Lead to oversimplified NN structures
 - >10% accuracy lost
- Remote Inference
 - **Compress** raw data before transmission
 - Limited **data compressibility** when accuracy loss is minimum.
- NN offloading
 - Use a local NN to sparsity & compress data
 - Higher compressibility but expensive local NNs
 - Extra transmission cost

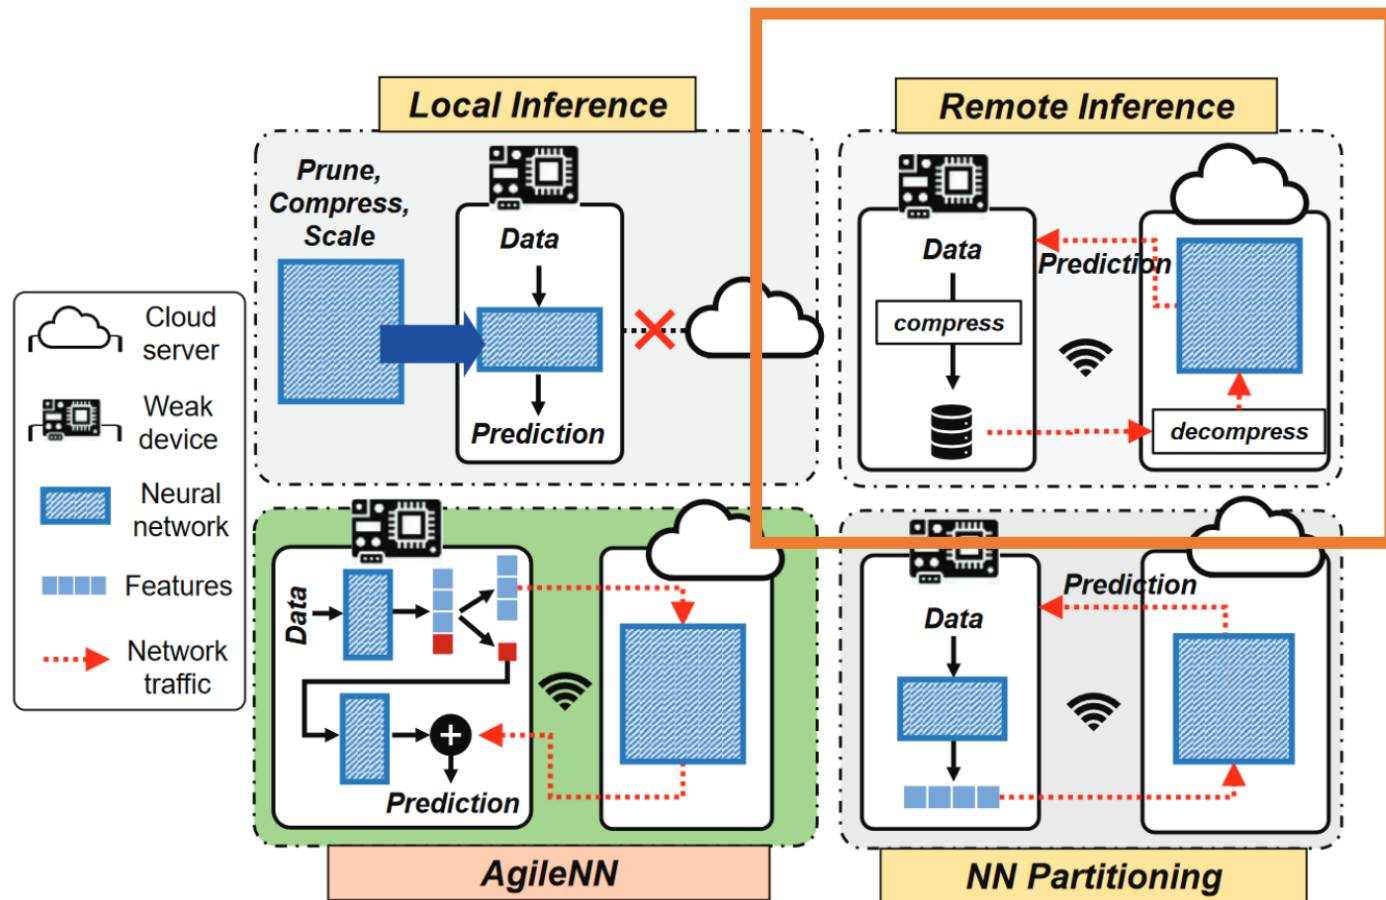


Figure 1: Existing work vs. AgileNN

Existing Work

- Local Inference
 - Prune, Compress, Network Architecture Search
 - Lead to oversimplified NN structures
 - >10% accuracy lost
- Remote Inference
 - Compress raw data before transmission
 - Limited data compressibility when accuracy loss is minimum.
- NN offloading
 - Use a local NN to sparsity & compress data
 - Higher compressibility but **expensive local NNs**

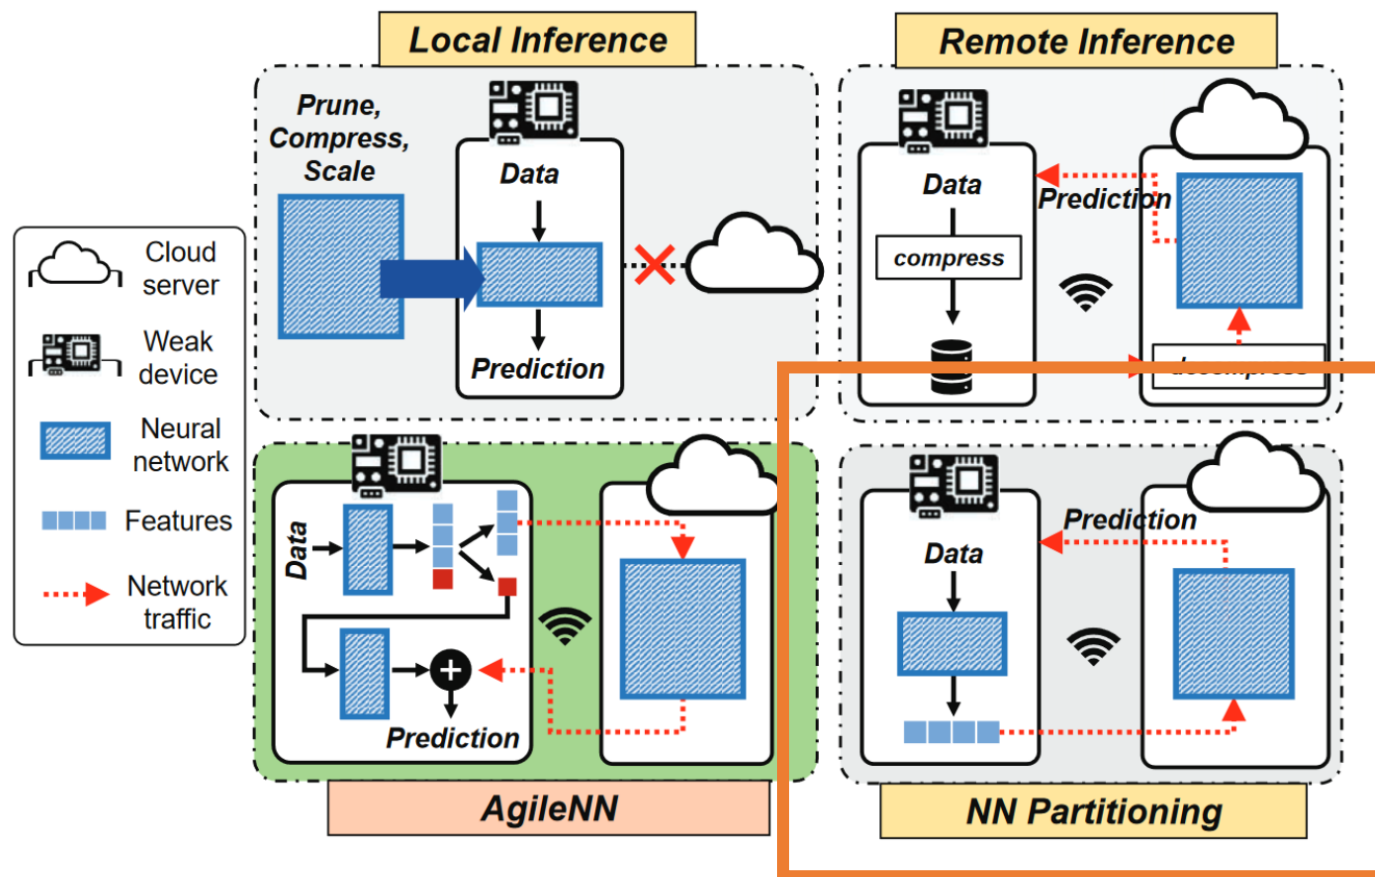



Figure 1: Existing work vs. AgileNN

AgileNN : From fixed to **data-centric** and **agile**



Data-centric





Incorporate the knowledge about different input data's **heterogeneity** in training



Agile Offloading

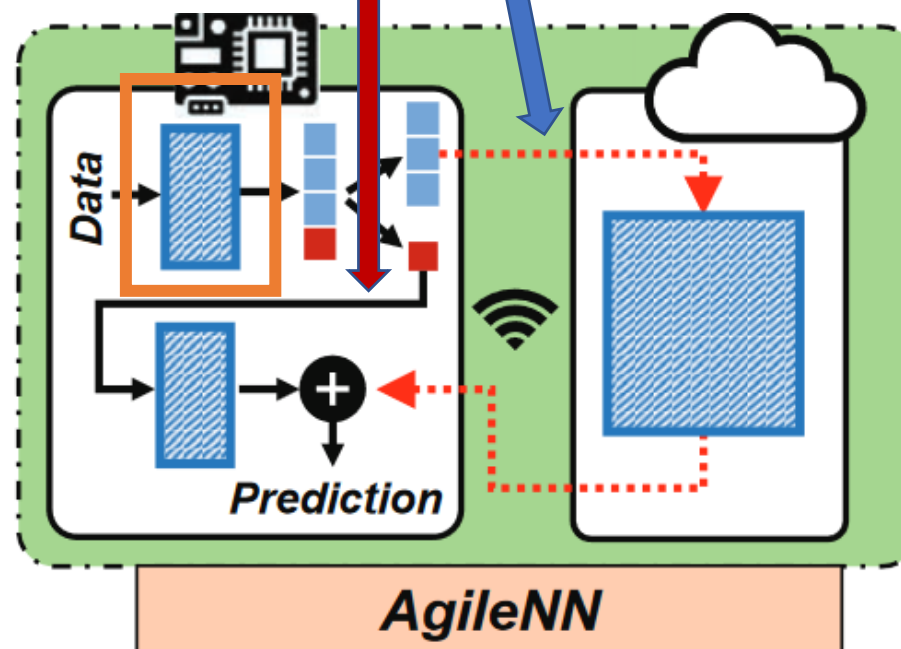
Adaptive partition to minimize the offloading cost

AgileNN : From fixed to **data-centric** and **agile**



Feature
Extractor

For every Input data:
TopK important features – local NN
Less important features – remote NN

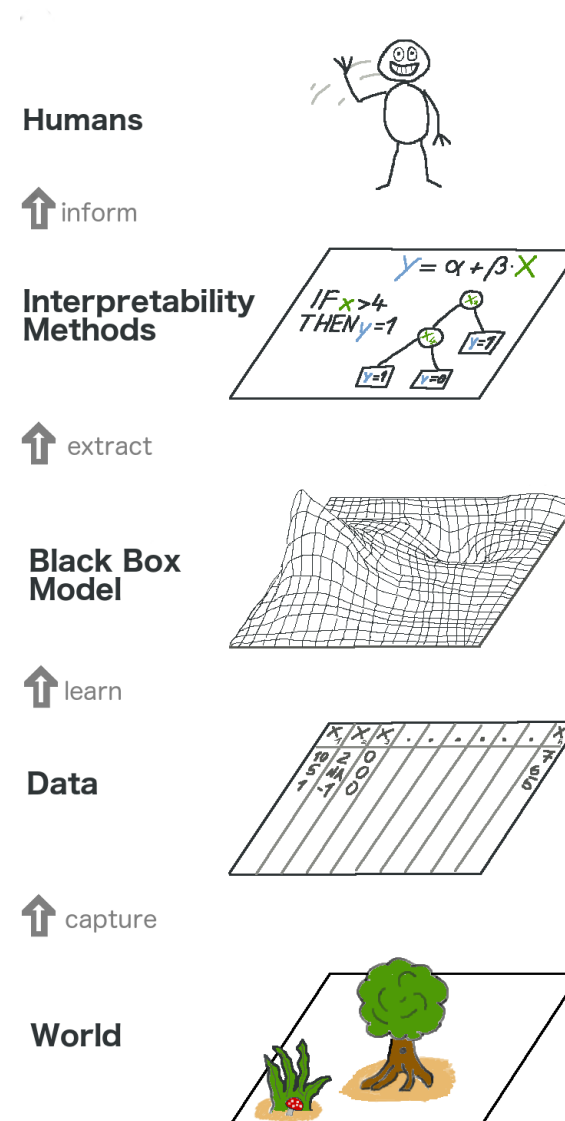


A Summary about Interpretable ML

- **Interpretability**
 - the degree to which a human can understand the cause of a decision.



Why it's a dog?



XAI Tools - Integrated Gradients

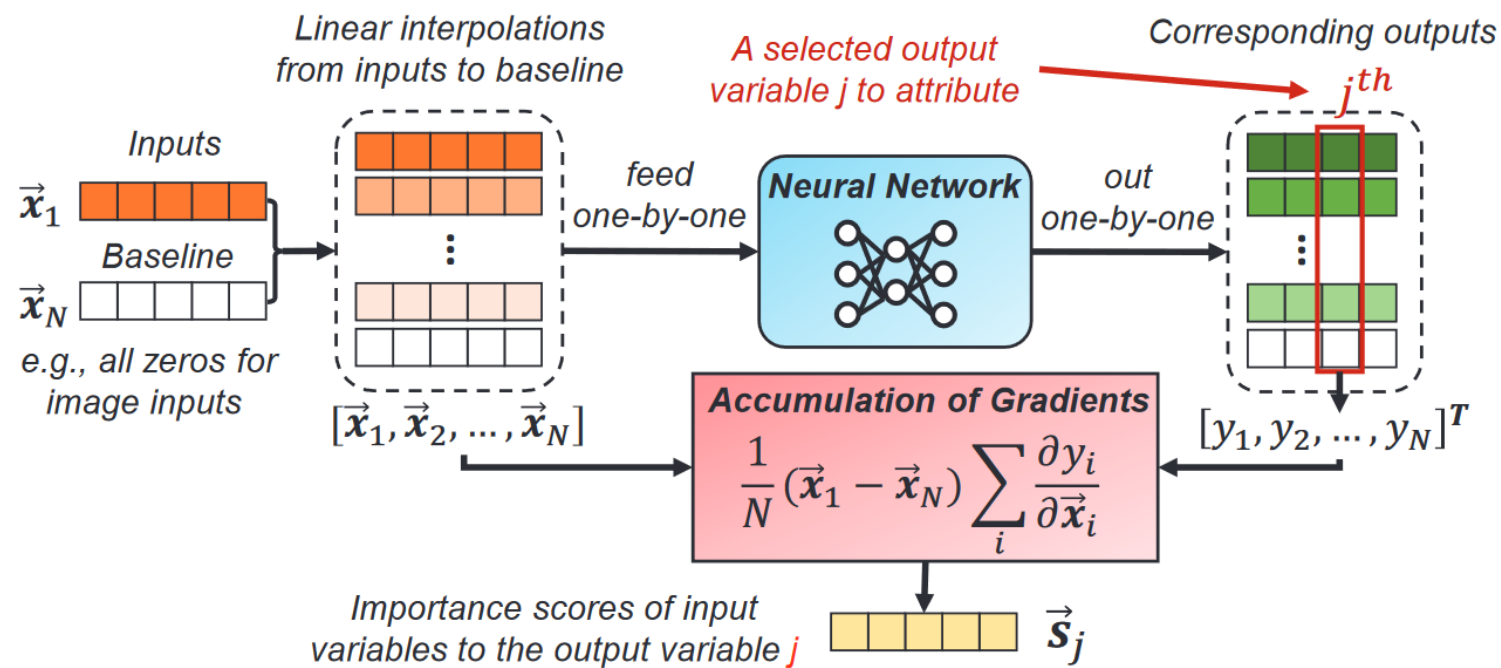


Figure 3: Integrated Gradients

Integrated Gradients aims to explain the relationship between the **model predictions** and its **features**.

XAI Tools - Integrated Gradients

Original image



Top label and score

Top label: reflex camera
Score: 0.993755

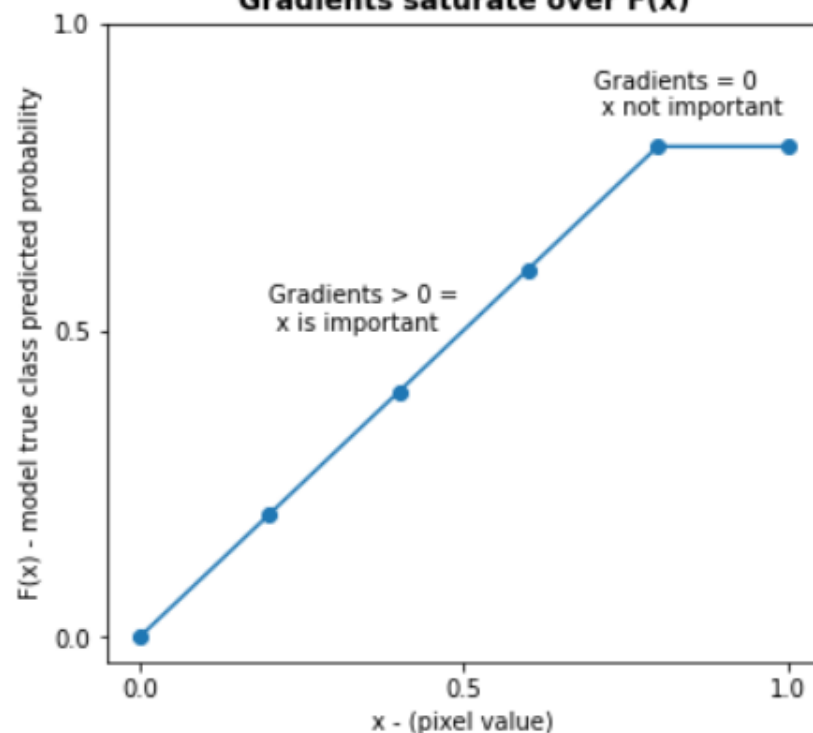
Integrated gradients



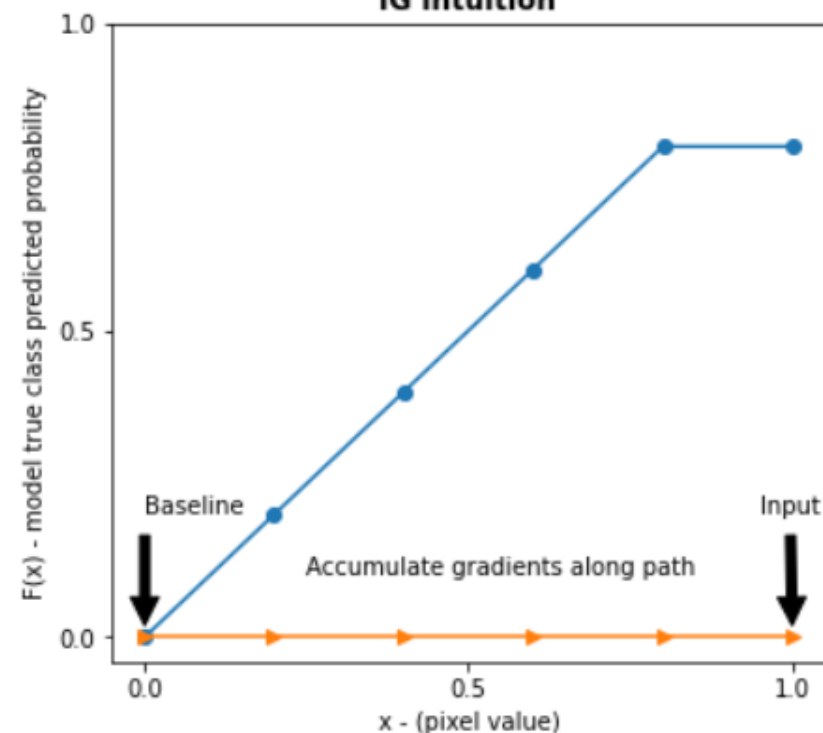
Gradients at image



Gradients saturate over $F(x)$



IG intuition



XAI Tools - Integrated Gradients

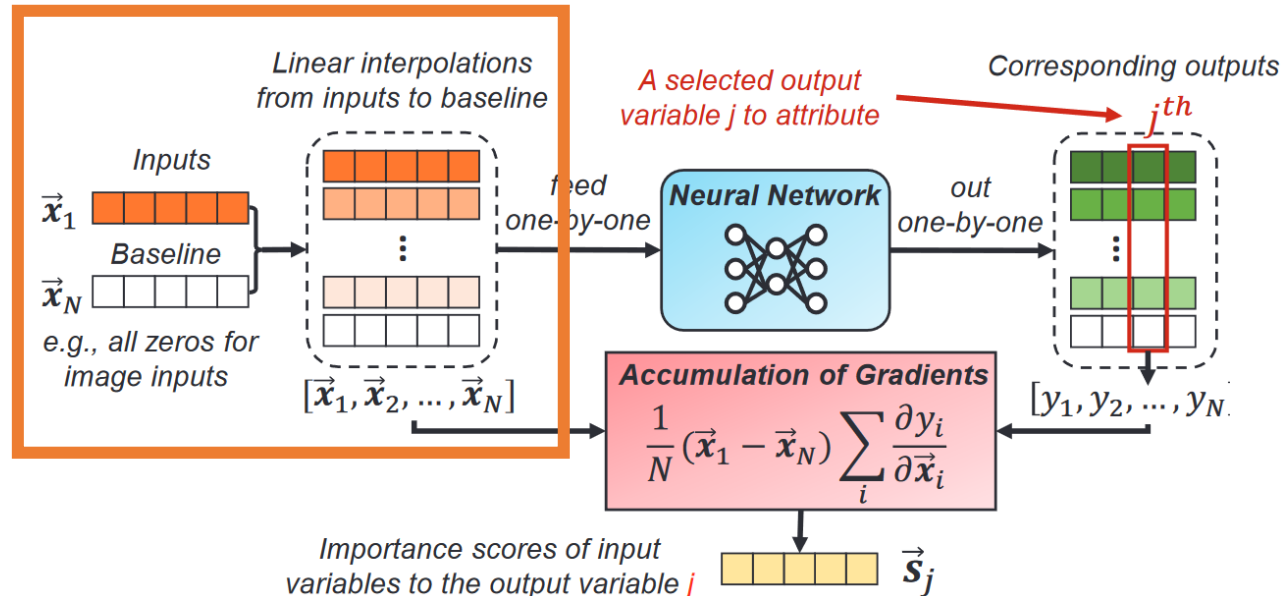



Figure 3: Integrated Gradients



1. Generate alphas α

2. Generate interpolated images = $(x' + \frac{k}{m} \times (x - x'))$



Integrated Gradients aims to explain the relationship between the **model predictions** and its **features**.

XAI Tools - Integrated Gradients

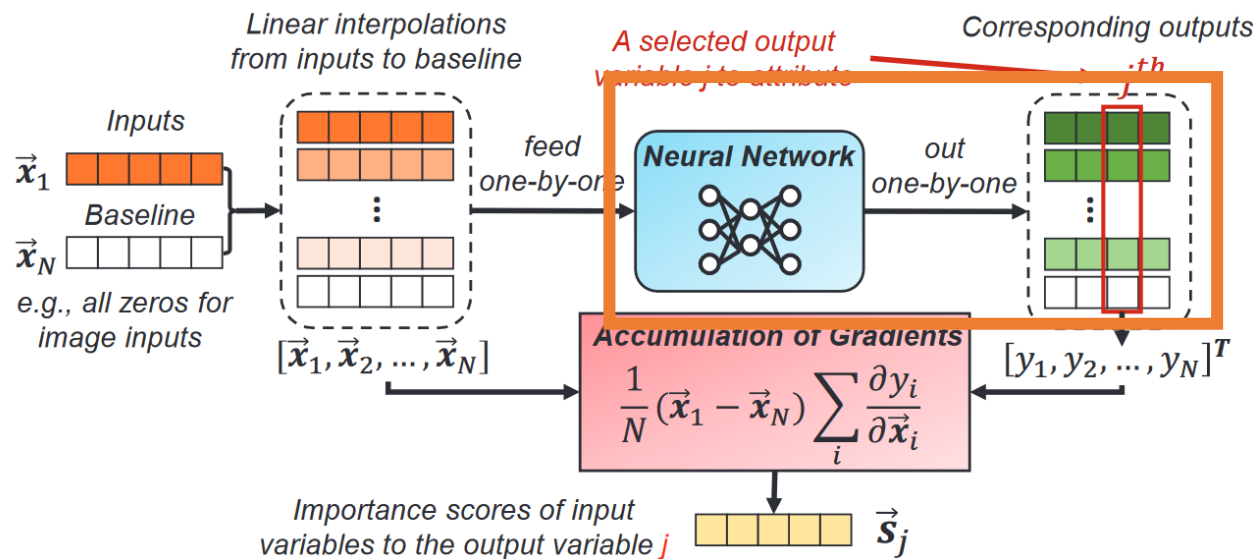
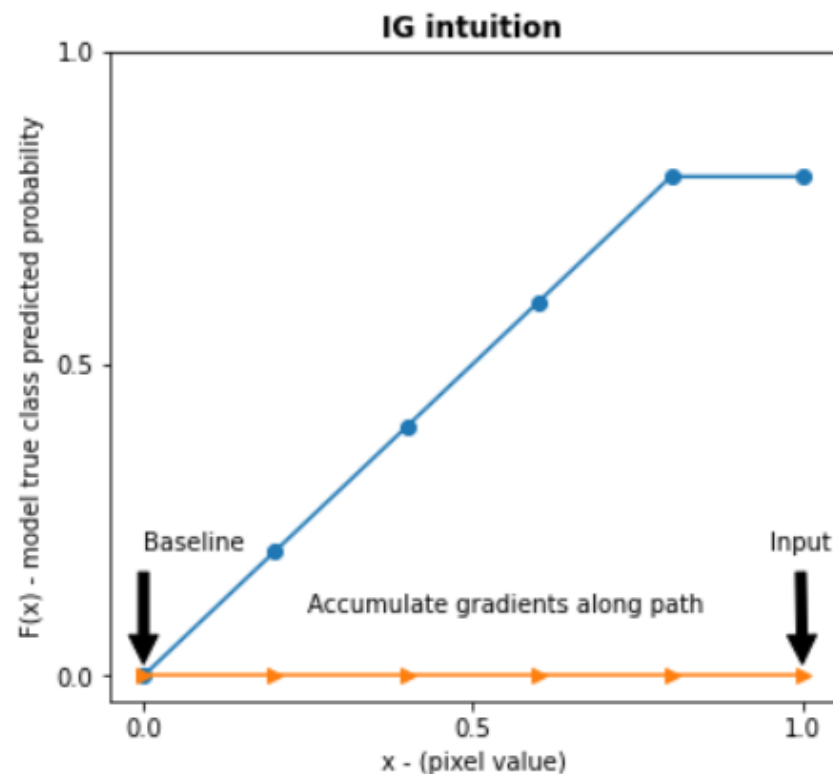


Figure 3: Integrated Gradients



3. Compute gradients between model F output predictions with respect to input features = $\frac{\partial F(\text{interpolated path inputs})}{\partial x_i}$

XAI Tools - Integrated Gradients

$$IntegratedGrads_i^{approx}(x) ::= (x_i - x'_i) \times \overbrace{\sum_{k=1}^m}^{\text{Sum } m \text{ local gradients}} \text{gradients}(\text{interpolated images}) \times \overbrace{\frac{1}{m}}^{\text{Divide by } m \text{ steps}}$$

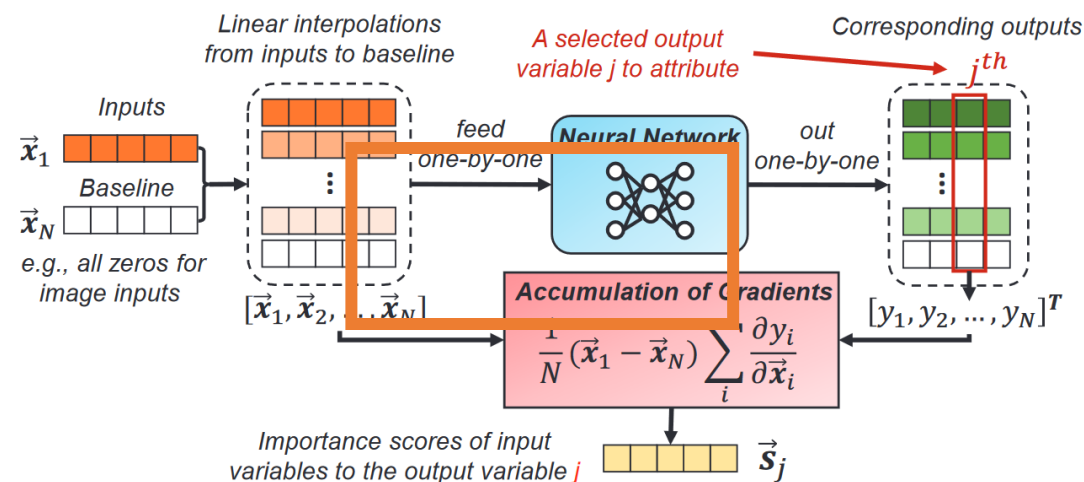


Figure 3: Integrated Gradients

XAI Tools - Integrated Gradients

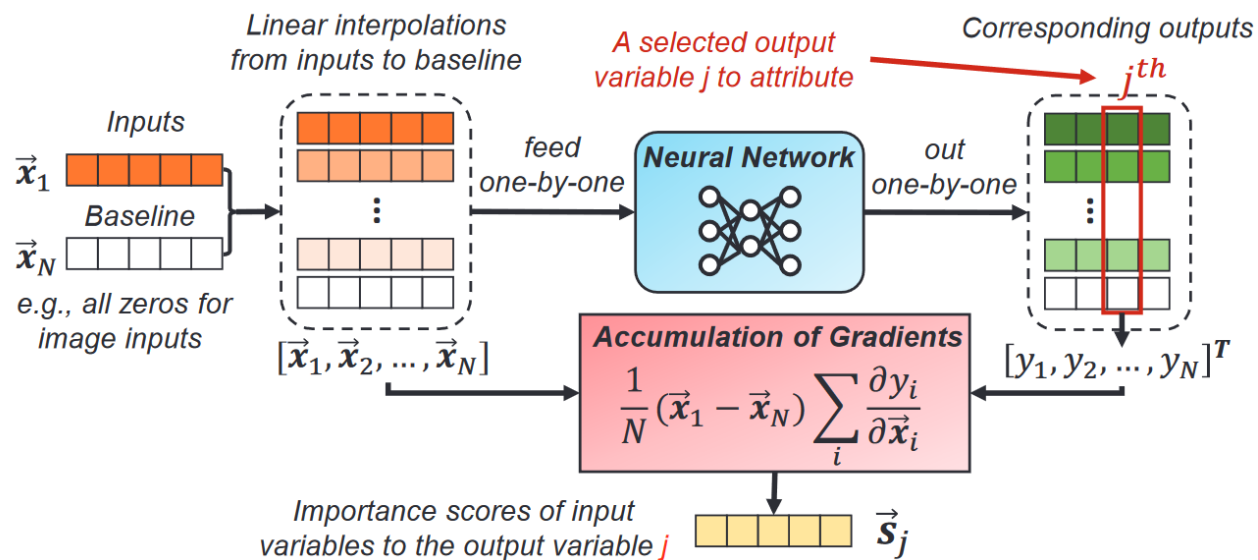
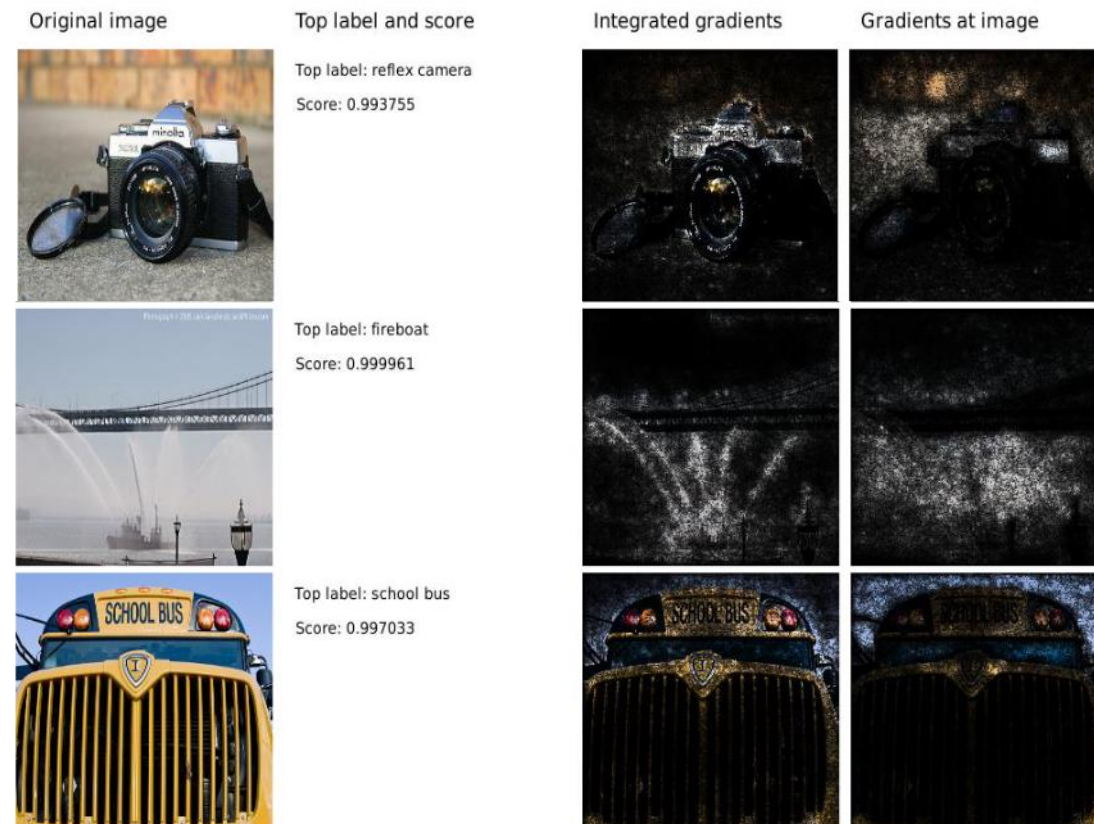



Figure 3: Integrated Gradients



Integrated Gradients aims to explain the relationship between the **model predictions** and its **features**.


XAI-enabled feature extractor

A typical XAI workflow in NN inference

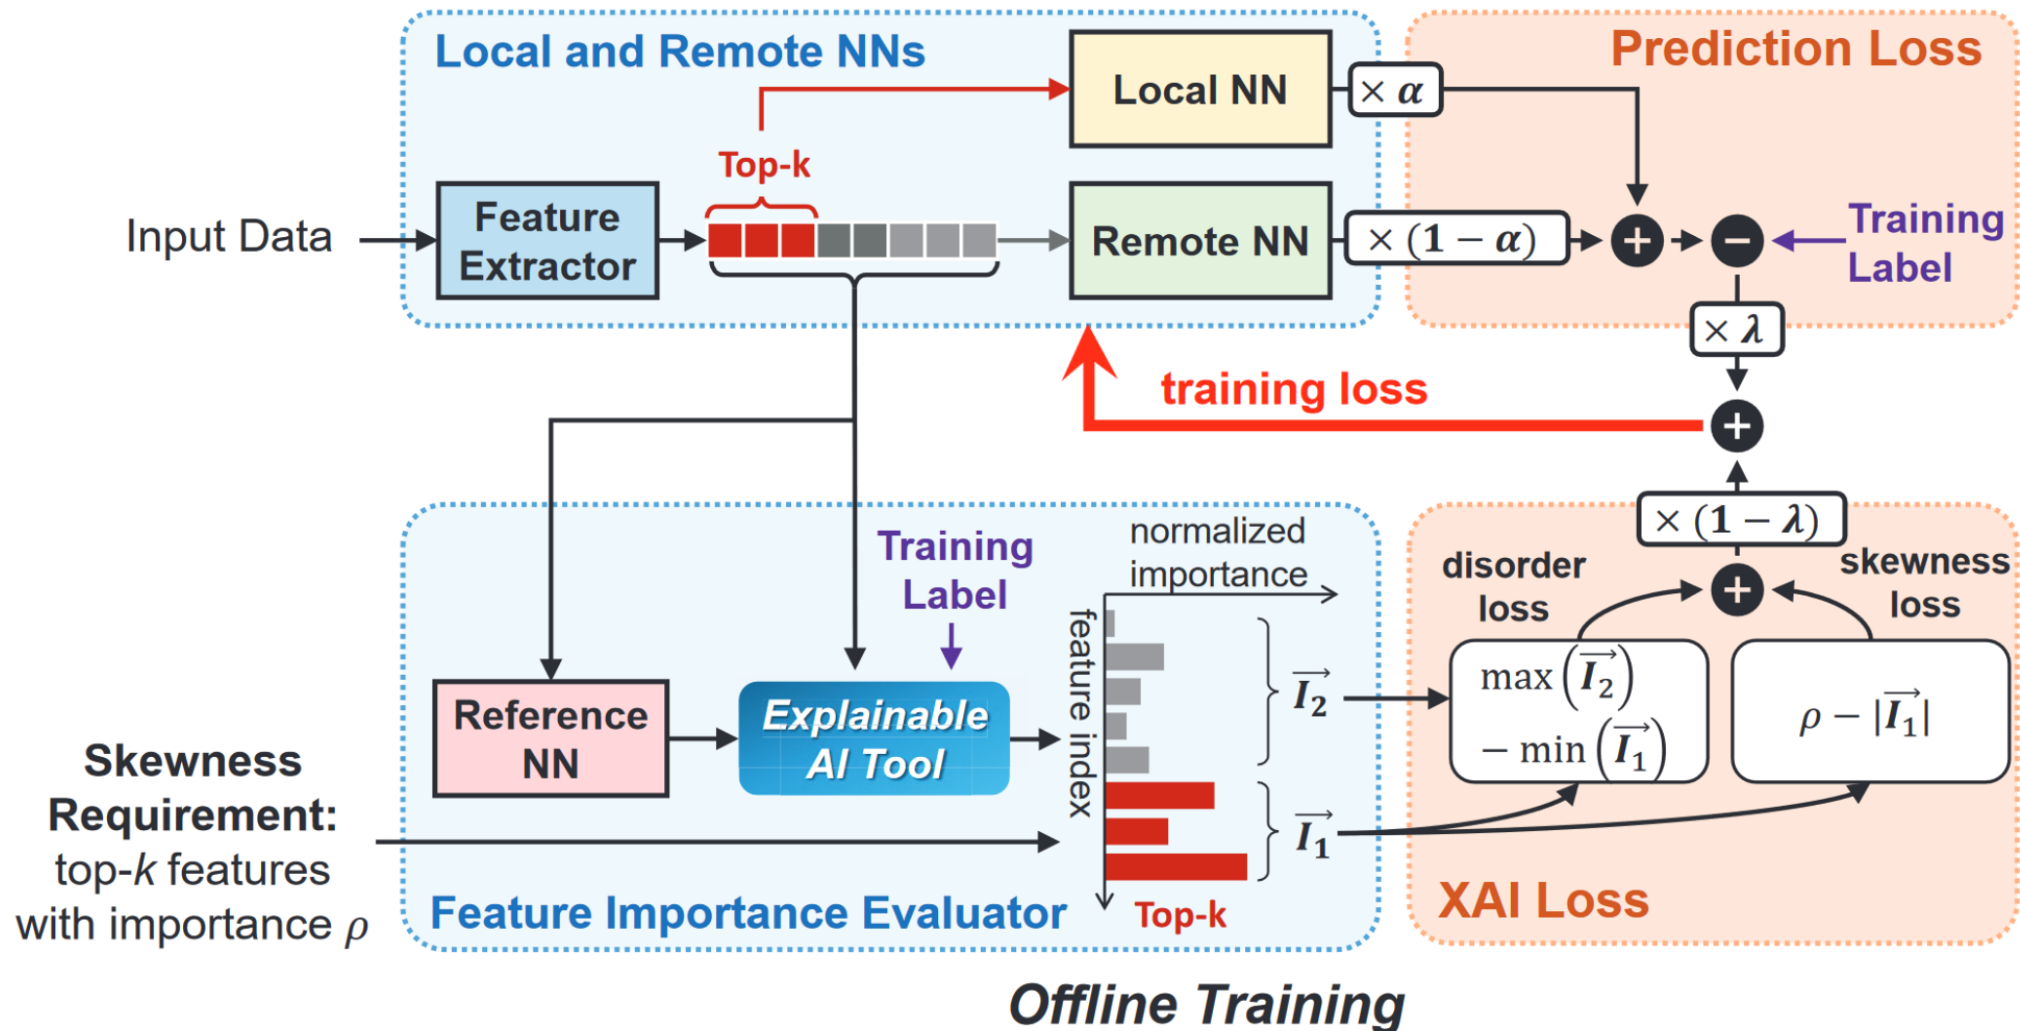


But Expensive!!

Instead, we suggest to address this challenge via agile NN offloading, which migrates the required computations in NN offloading from **online inference** to **offline learning**.



AgileNN : Offline Training



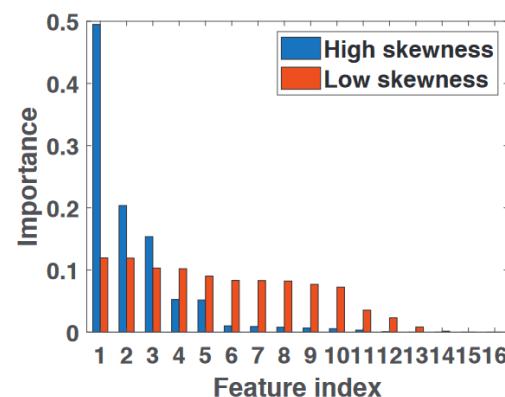


XAI Tools - Integrated Gradients

- Consider 1

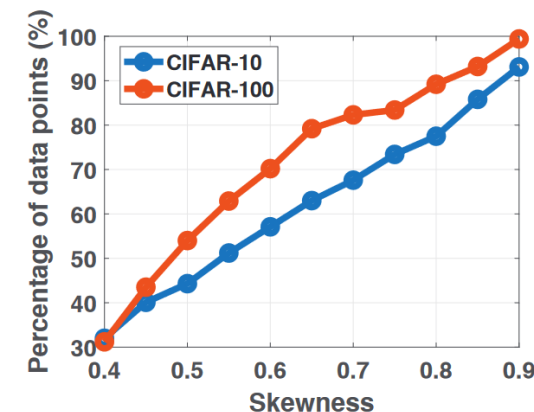
Skewness may not always exist in every input data.

- Consider 2

The accuracy of feature importance evaluation builds on accurate NN inference in advance.



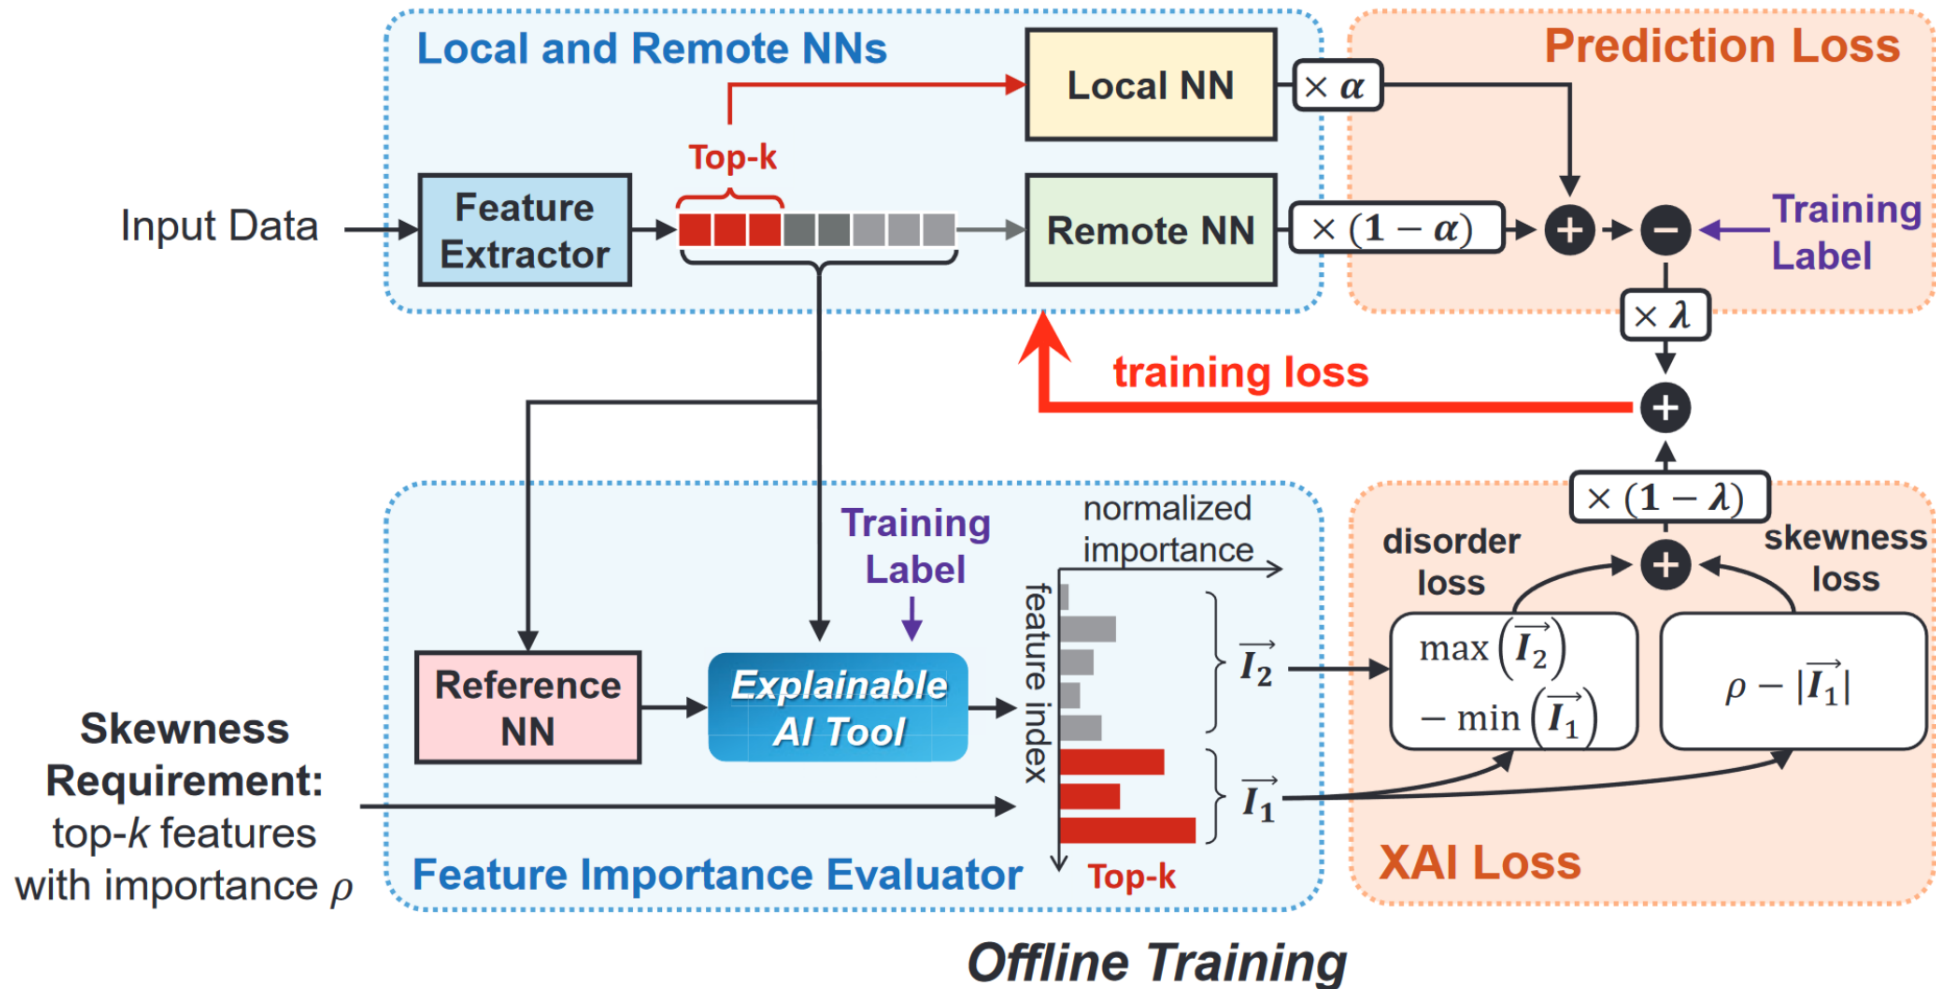
(a) Different levels of skewness



(b) Skewness CDF

Figure 4: Skewness of feature importance. Skewness is measured as the normalized importance of the top 20% features, using the MobileNetV2 model [55].

AgileNN : Offline Training



Enforce skewed distribution of features

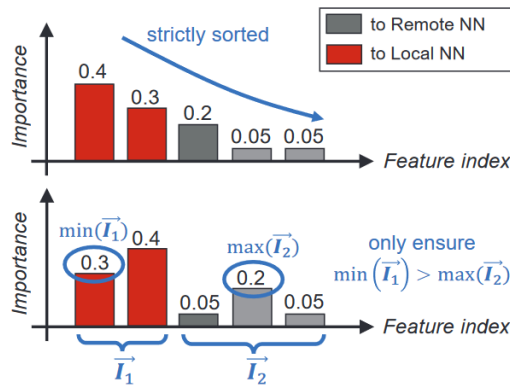
Disorder loss

Ensure topmost important features are extracted into **topK channels**


$$L_{\text{disorder}} = \max \left(0, \max(\vec{I}_2) - \min(\vec{I}_1) \right)$$

Skewness Loss

Enhance the important of topK features to **ensure compressibility** of the others



(a) Different feature orders



(b) Effectiveness of reordering




Figure 9: Feature reordering

Enforce skewed distribution of features


Disorder loss

Ensure topmost important features are extracted into **topK channels**


Skewness Loss

Enhance the important of topK features to **ensure compressibility** of the others

$$L_{\text{skewness}} = \max \left(0, \rho - |\vec{I}_1| \right)$$



AgileNN : Offline Training



Combining Local and Remote Predictions

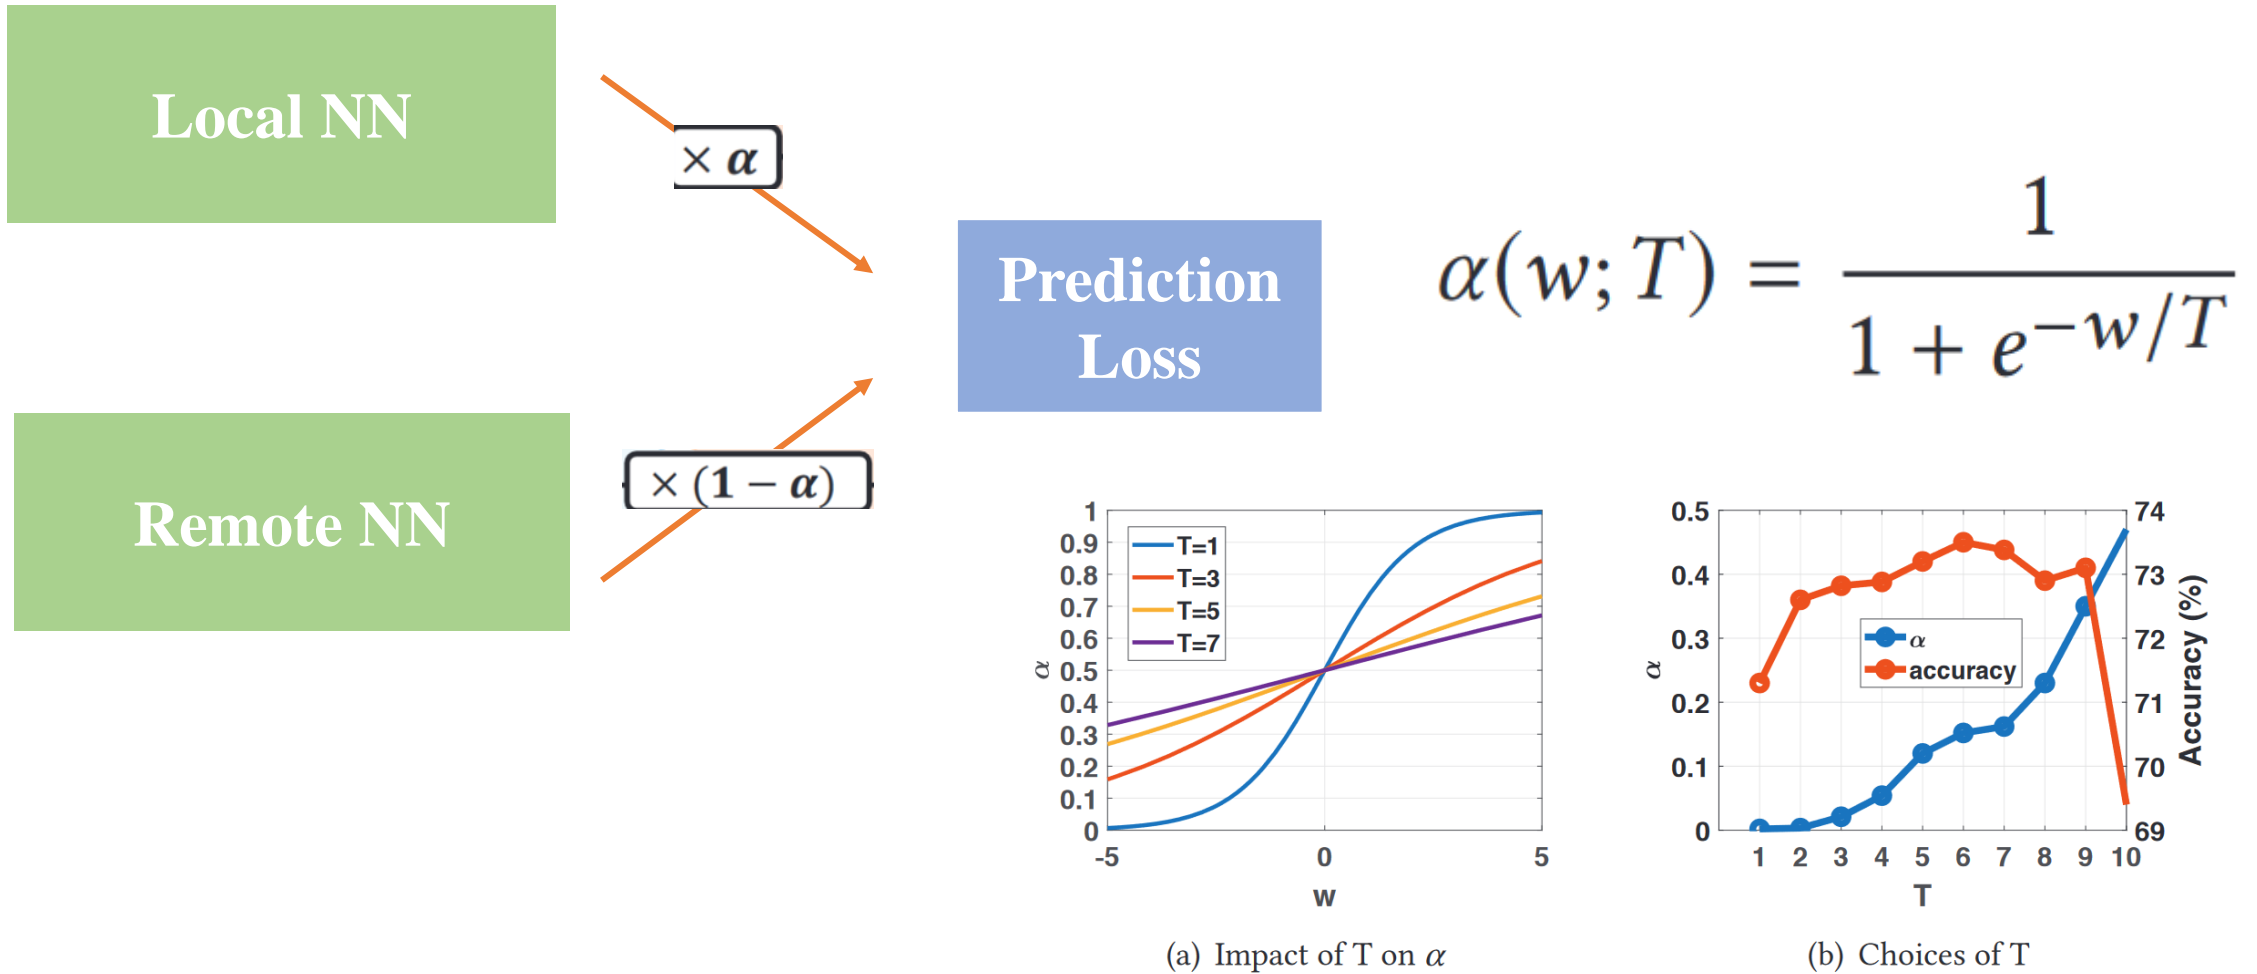


Figure 8: Prediction weighting with α

Combined Training Loss

Disorder loss

Ensure topmost important features are extracted into topK channels

Skewness Loss

Enhance the important of topK features to ensure compressibility of the others

$$L = \lambda \cdot L_{\text{prediction}} + (1 - \lambda) \cdot (L_{\text{skewnss}} + L_{\text{disorder}})$$

Local NN

$\times \alpha$

Remote NN

$\times (1 - \alpha)$

Prediction
Loss




Figure 10: Impact of λ on CIFAR-100 dataset

Pre-processing the feature extractor



Figure 6: Training stability with different numbers of convolutional layers in feature extractor

Pre-processing the feature extractor

Algorithm 1 Selecting the k initial feature channels

Input: $\mathbf{D}_{\text{train}}$: the training dataset with N samples;
 $\mathcal{T}_{\text{XAI}}(\cdot)$: XAI-enabled Feature Importance Evaluator;
 $\mathcal{E}(\cdot)$: Feature extractor that outputs C channels

Output: (j_1, j_2, \dots, j_k) : The k selected feature channels.

```
1:  $(p_1, p_2, \dots, p_C) \leftarrow 0$  // initialize
2: for each  $d_i \in \mathbf{D}_{\text{train}}$  do
3:    $F \leftarrow \mathcal{E}(d_i)$  // extract features
4:    $I \leftarrow \mathcal{T}_{\text{XAI}}(F)$  // evaluate feature importance
5:    $F_{\text{sorted}} \leftarrow \text{sort}_I(F)$  // sort features by their importance in
   descending order
6:    $F_{\text{top-}k} \leftarrow F_{\text{sorted}}[1 : k]$  // extract the top- $k$  features with high
   importance
7:   for  $c = 1, \dots, C$  do
8:     if  $F[c] \in F_{\text{top-}k}$  then
9:        $p_c \leftarrow p_c + 1/N$ 
10:  $R \leftarrow \text{argsort}(p_1, p_2, \dots, p_C)$  // get the ranking of channels
   by their likelihood
11:  $(j_1, j_2, \dots, j_k) \leftarrow R[1 : k]$  // decide top- $k$  channels
```




Figure 11: Effectiveness of Pre-processing

Pre-processing the feature extractor




Figure 7: Pre-processing the feature extractor




Figure 12: Training the mapping layer

Rearrange the indices to reduce unsorted cases in training data

XAI Tools - Integrated Gradients

- Consider 1

Skewness may not always exist in every input data.

- Consider 2

The accuracy of feature importance evaluation builds on accurate NN inference in advance.

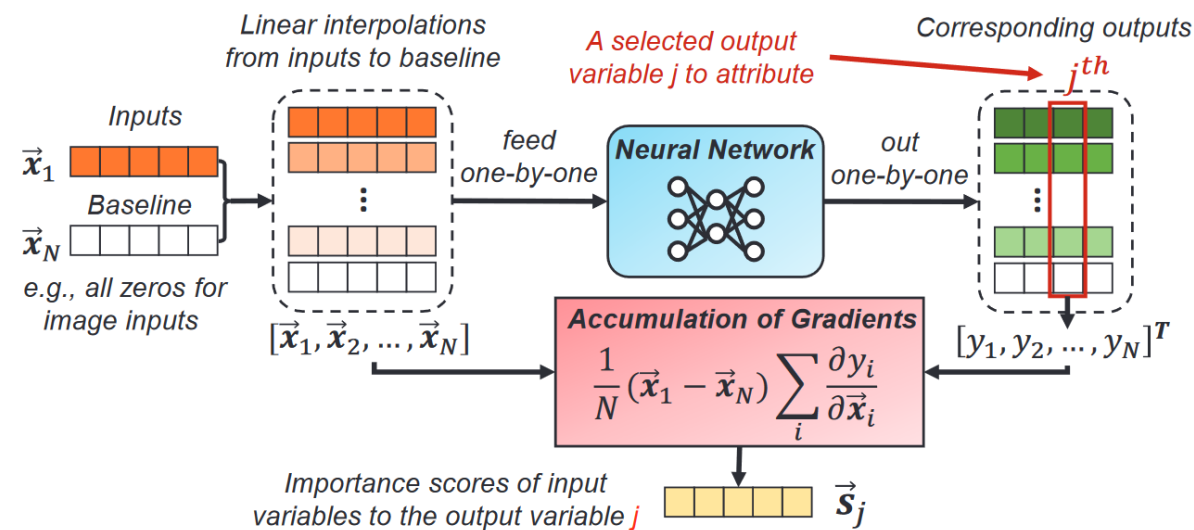
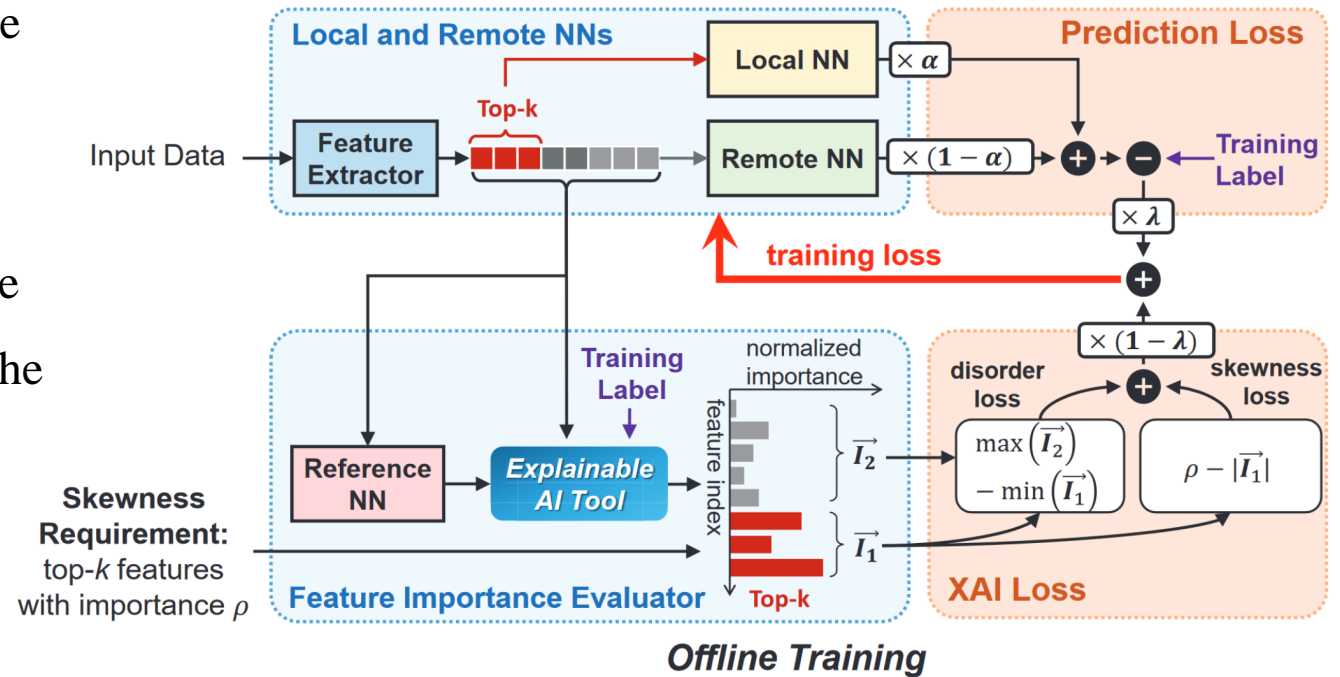


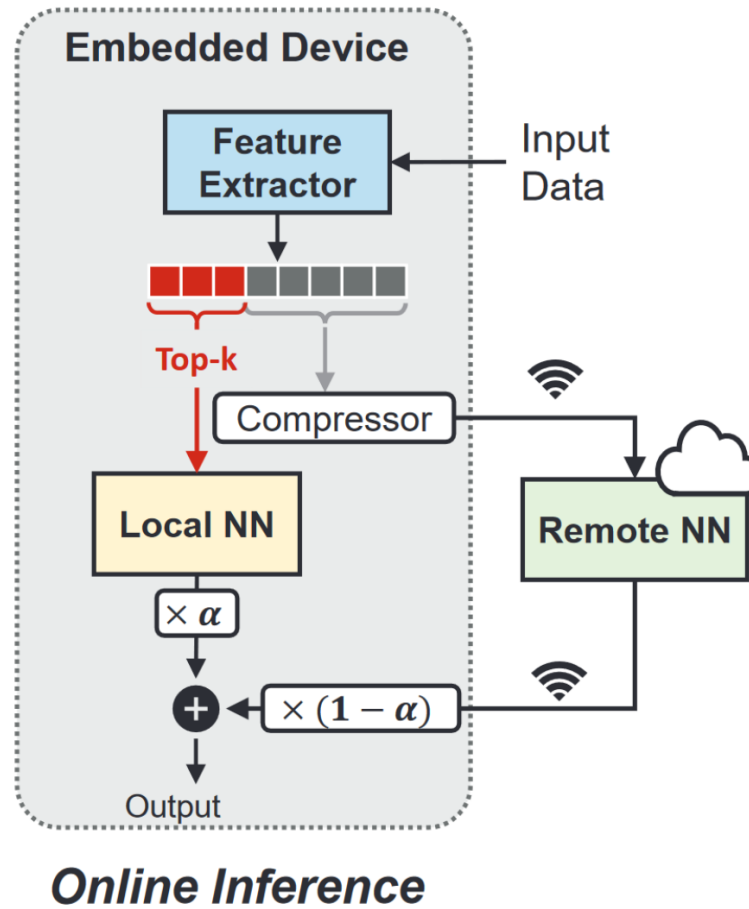
Figure 3: Integrated Gradients

AgileNN : Offline Training

- The **higher** the skewness is (i.e., smaller k and larger ρ), the **lower** resource consumption will be at the local.
- But the NN inference is **more affected** due to the feature extractor's non-linear transformation in the feature space.
- AgileNN allow **flexible tradeoffs** between the accuracy and cost of NN inference on embedded devices, without incurring any extra computing or storage cost.



AgileNN : Online Inference



Implementation

- Local device

STM32F746 board, 216MHz, 320kb SRAM, 1MB FRAM

ESP-WROOM-02D WiFi module @ 6Mbps

- Remote device

Dell-Precision 7820 workstation




Figure 13: Devices in our implementation

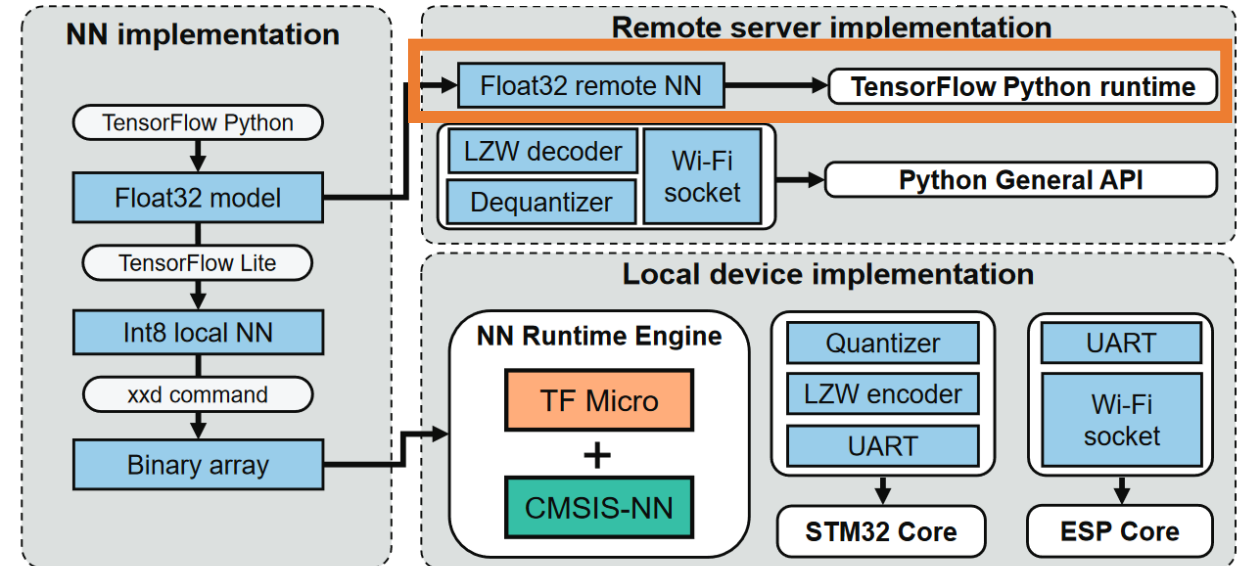
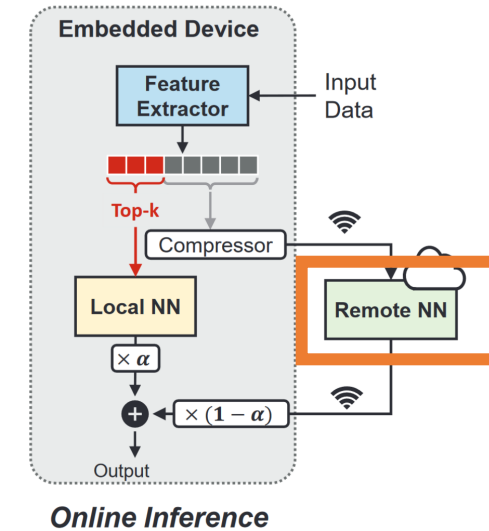


Figure 14: AgileNN implementation

Implementation

- Local device

STM32F746 board, 216MHz, 320kb SRAM, 1MB FRAM

ESP-WROOM-02D WiFi module @ 6Mbps

- Remote device

Dell-Precision 7820 workstation




Figure 13: Devices in our implementation

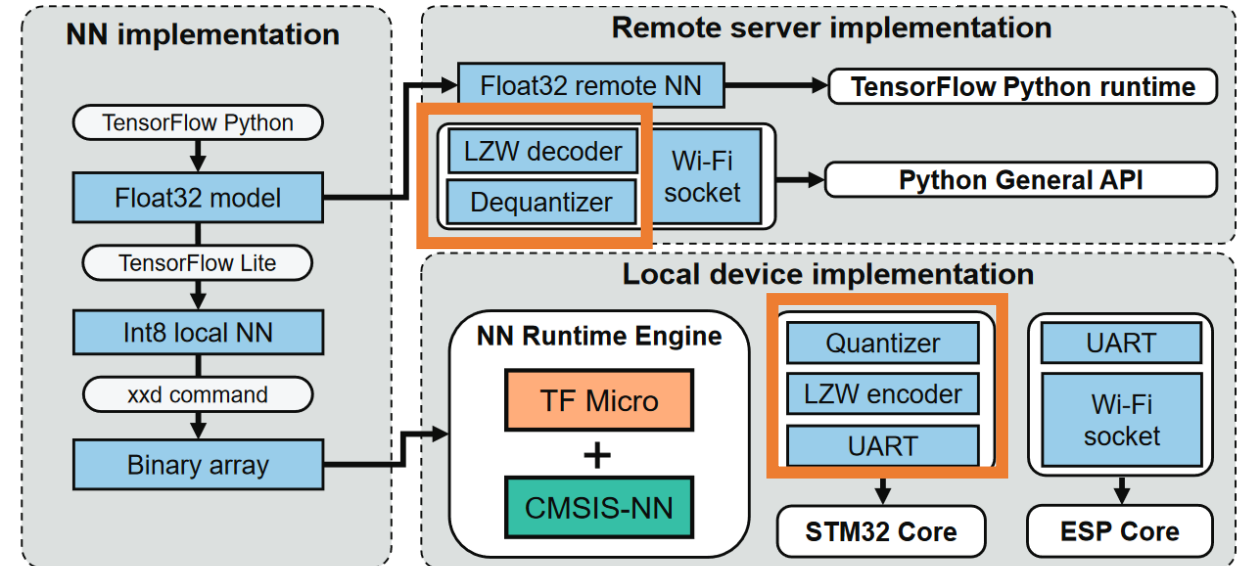
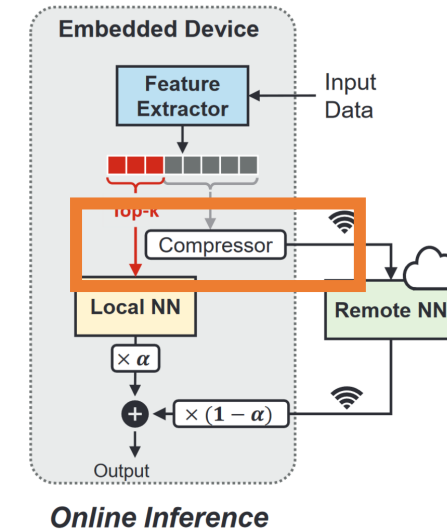


Figure 14: AgileNN implementation

Implementation

- Local device

STM32F746 board, 216MHz, 320kb SRAM, 1MB FRAM

ESP-WROOM-02D WiFi module @ 6Mbps

- Remote device

Dell-Precision 7820 workstation




Figure 13: Devices in our implementation

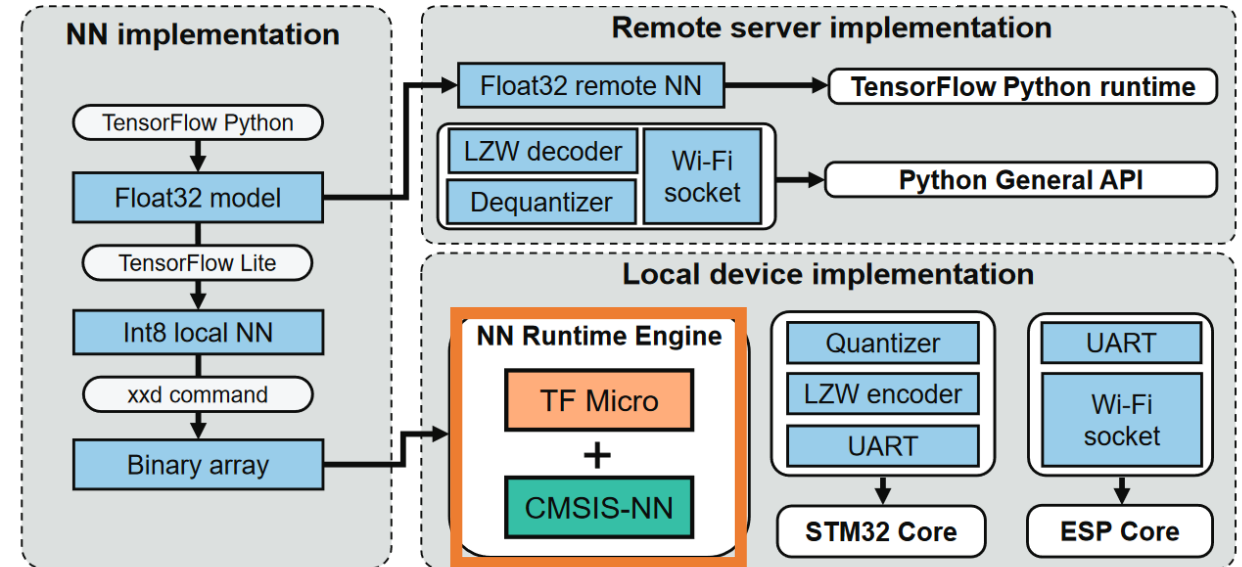
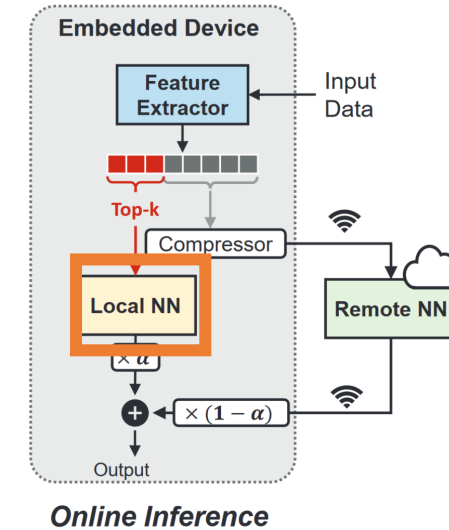


Figure 14: AgileNN implementation

Implementation

- Local device

STM32F746 board, 216MHz, 320kb SRAM, 1MB FRAM

ESP-WROOM-02D WiFi module @ 6Mbps

- Remote device

Dell-Precision 7820 workstation




Figure 13: Devices in our implementation

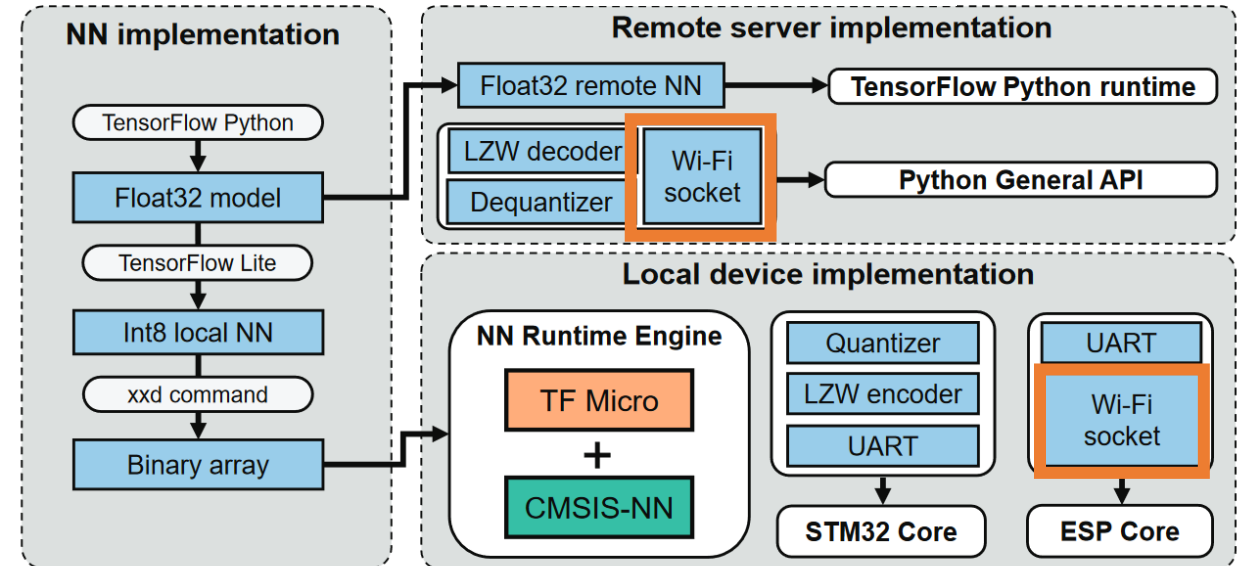
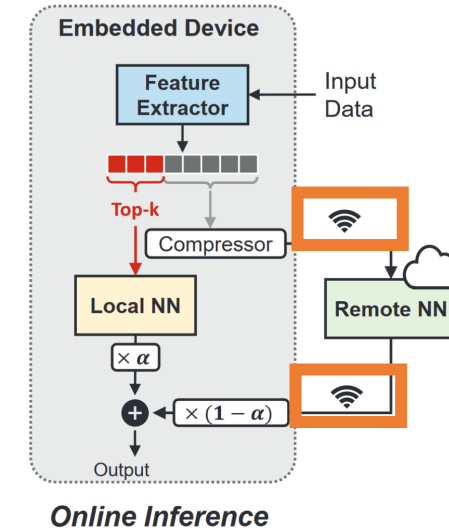


Figure 14: AgileNN implementation

Evaluation Setup

- **Baselines**

- MCUNet [1] – NAS to find the best local NN structures – **Local NN**
- DeepCOD [2] – use a NN-based encoder – **Remote NN**
- SPINN [3] – early-exit structures NN – **NN offloading**
- Edge-only Inference – compress and offload raw data

- **Datasets**

- **CIFAR-10/100**
- **SVHN**
- **ImageNet-200**

- **Reference NN**


- **EfficientNetV2 CNN**

[1] Lin, Ji, et al. "Mccunet: Tiny deep learning on iot devices." *Advances in Neural Information Processing Systems* 33 (2020): 11711-11722.


[2] Yao, Shuochao, et al. "Deep Compressive Offloading: Speeding Up Edge Offloading for AI Services." *GetMobile: Mobile Computing and Communications* 25.1 (2021): 39-42.

[3] Laskaridis, Stefanos, et al. "SPINN: synergistic progressive inference of neural networks over device and cloud." *Proceedings of the 26th annual international conference on mobile computing and networking*. 2020.

Evaluation - Training Convergence and Cost



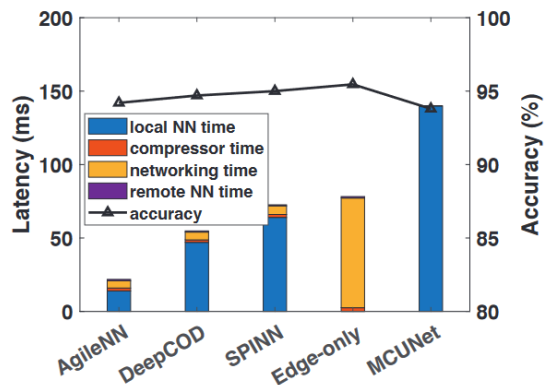
(a) Test accuracy




(b) Test loss

Evaluation - Accuracy and Latency

- *AgileNN* reduces end-to-end latency by 2x – 2.5x




(a) CIFAR-10




(b) CIFAR-100

Dataset	CIFAR-10	CIFAR-100	SVHN	ImageNet
Reduction	43.7%	15.8%	72.3%	20.8%

Table 2: Reduction of transmitted data size, compared to DeepCOD [65]




(c) SVHN




(d) ImageNet-200

Evaluation - **Accuracy** and Latency

- *AgileNN* reduces end-to-end latency by 2x – 2.5x



(a) CIFAR-100



(b) SVHN

Figure 17: Accuracy with different compression rates

Evaluation – Different System Settings

- **Impact of Local CPU Frequency**

- 64MHz – 216MHz
- Reduces latency by 2.1x to 2.5x




Figure 22: The impact of different CPU Frequencies

- **Impact of Network Bandwidth**

- Bluetooth 270kbps, 2Mbps
- WiFi 6Mbps





Figure 23: The impact of different wireless bandwidths


Evaluation – Local Resources Consumption

- **Local energy consumption**

- 1.6x – 2.5x more efficient than DeepCOD
- 8x more efficient than MCUNet



(a) CIFAR-100




(b) ImageNet-200


Figure 19: Local energy consumption per NN inference run

- **Memory and storage usage**

- SRAM to Memory
- FRAM to Storage
- Higher accuracy, and save 40%-50% memory and >50% storage



(a) CIFAR-100



(b) ImageNet-200

Figure 20: Memory and storage usage

Evaluation – Effectiveness of Skewness Manipulation

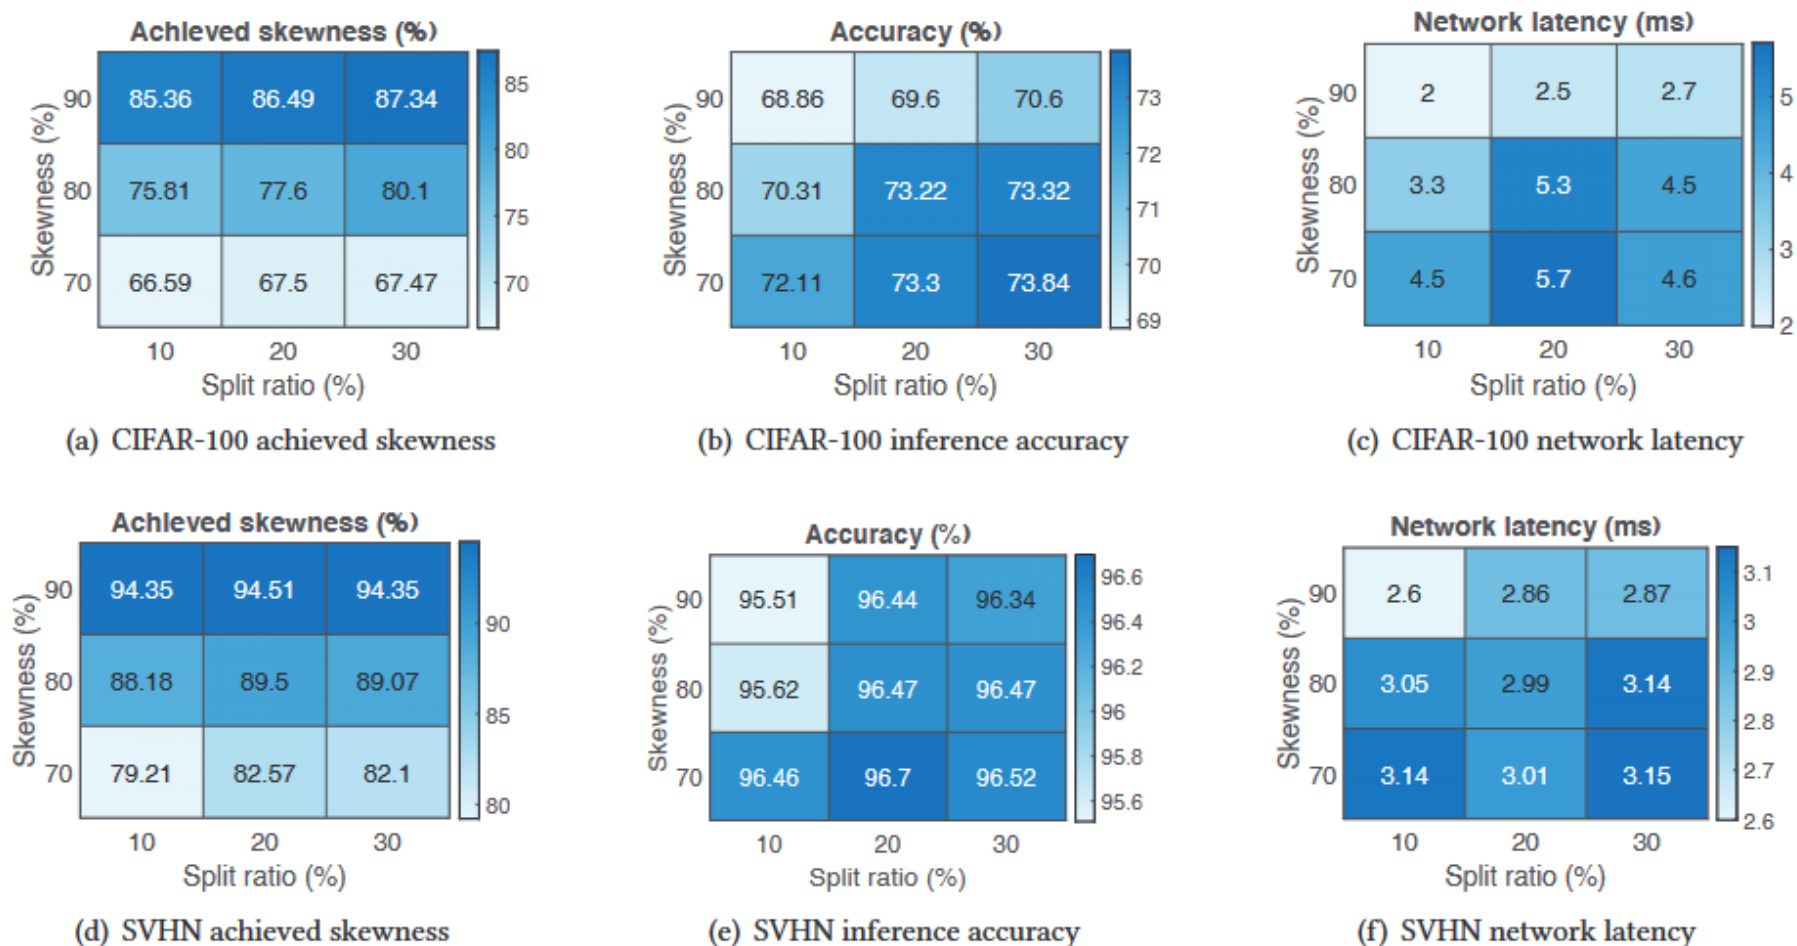


Figure 21: Effectiveness of skewness manipulation with different requirements of feature importance skewness

Evaluation – Different System Settings

- **Impact of Local CPU Frequency**

- 64MHz – 216MHz
- Reduces latency by 2.1x to 2.5x




Figure 22: The impact of different CPU Frequencies

- **Impact of Network Bandwidth**

- Bluetooth 270kbps, 2Mbps
- WiFi 6Mbps

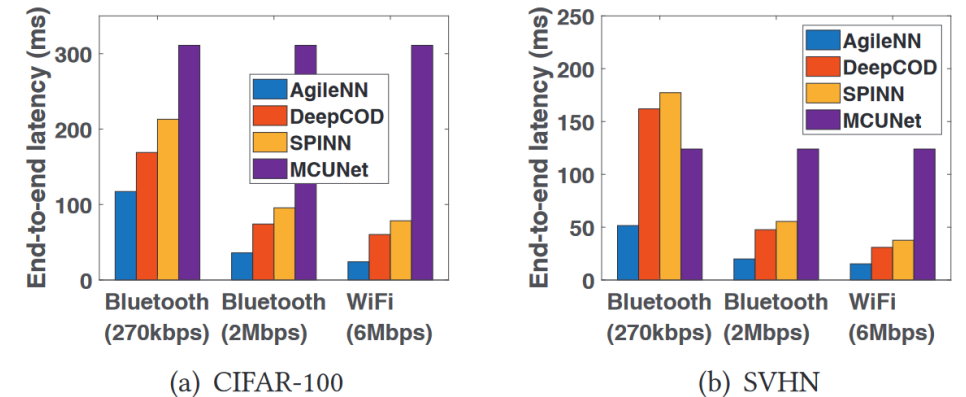


Figure 23: The impact of different wireless bandwidths

Evaluation - Choices of XAI techniques

- **Gradient Saliency**
 - evaluate each feature's importance by injecting noise.
- **Integrated Gradients**
 - It aggregates more interpolations of NN outputs' gradients
 - IG is more computationally.




Figure 24: Different XAI techniques

Evaluation - Choices of XAI techniques

- **Gradient Saliency**
 - evaluate each feature's importance by injecting noise.
- **Integrated Gradients**
 - It aggregates more interpolations of NN outputs' gradients
 - IG is more computationally.




Figure 24: Different XAI techniques

Advantage

- First to integrate **XAI techniques** in to NN offloading systems
- Shifts the rationale of offloading from fixed to **data-centric & agile**
- **>6x** lower latency and **>8x** resources consumption for **extremely weak** devices

Disadvantage

- XAI tools do not consider the **relationship** between features.
- Hard for **static NN** models to adopt to new data and different application scenarios.