

BeeHive: Sub-second Elasticity for Web Services with Semi-FaaS Execution

Edge System Reading Group @ SEU

Haodong Tian

June 1, 2023.

Introduction – FaaS

- Real-world web environment stimulates demand for **resource elasticity**.
- Compared to others, FaaS provides:
 - Automatically scales applications in a finer granularity (functions)
 - Shorter reaction time
 - Pay-as-you-go model for cost-efficient computation

Introduction – FaaS

- Prior work has proposed to run various applications as FaaS functions, including video processing, software compilation, micro-services, etc.
- FaaS encounters challenges when deploying traditional **monolithic web applications**.

Motivation – Tackling Request Bursts with FaaS

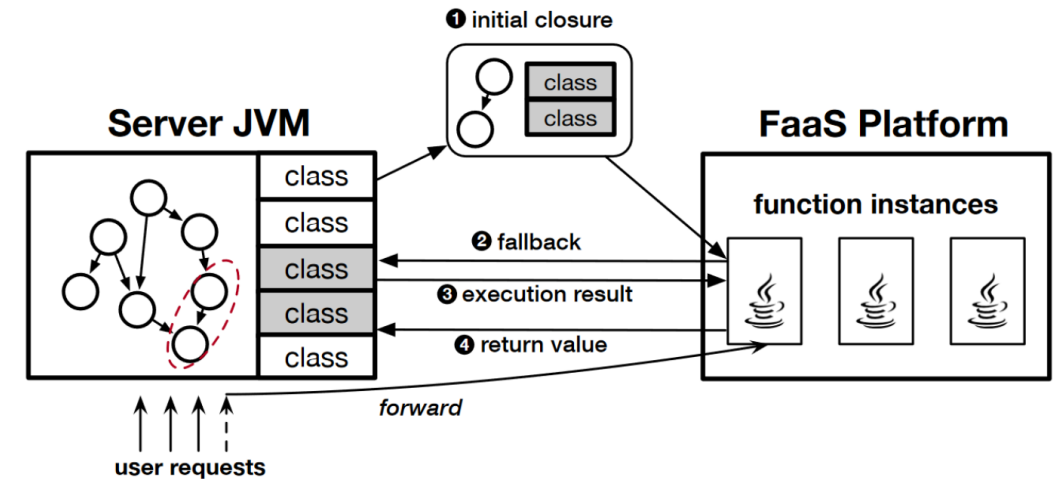
- Request bursts are long-term enemies for web applications.
- Cloud vendors (taking AWS as example) have provided various solutions for resource scaling:
 - Reserved Instance (RI) *high cost*
 - On-demand instance *long instance creation time*
 - Burstable instance *similar to RI but uses a different billing model*
 - Fargate *granularity and billing not so flexible*
 - Lambda (FaaS) *rapid, elastic, and automatic fashion*

Motivation – Initial Approaches of Applying FaaS to Web Applications

- It is not trivial to adapt existing web applications to FaaS platforms.
- Three different methods to migrate an enterprise-level web service into FaaS platform for execution:
 - Method 1: direct execution *violates stateless and lightweight assumptions*
 - Method 2: manual rewriting *too complicated*
 - Method 3: static code analysis *too dynamic to be statically analyzed*
- An execution model for web applications to leverage the power of FaaS should conform to the following principles:
 - *Partial. Automatic. Dynamic.*

System Design – Offloading-Based Semi-FaaS with BeeHive

- BeeHive mainly contains two parts: **long-running servers** and **FaaS platforms**.
- When facing request bursts, BeeHive controls servers to offload a part of its workload as functions to FaaS platforms for execution, while the rest is still handled by the server (namely **Semi-FaaS**).



System Design – Reducing the Performance Overhead of Fallbacks

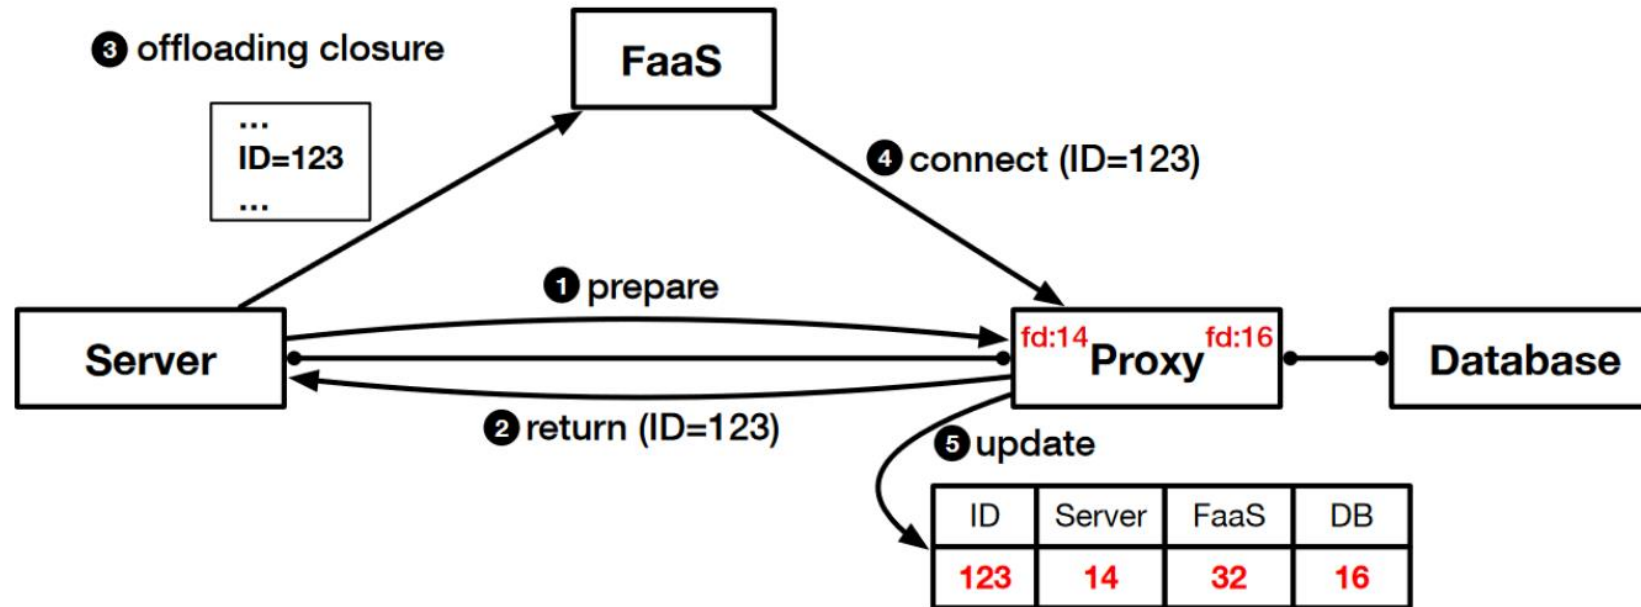
- **Native invocations**
- In web applications, **native invocations** are heavily used (System.arraycopy, Thread.currentThread, etc) but treated as not offloadable. Returning to servers for handling would cause prohibitive overhead.
- Native methods could be divided into four categories:
 - Pure on-heap operations. *can be directly executed on FaaS*
 - Hidden states. *marshal the hidden states into closure*
 - Network-related. *will be discussed later...*
 - Stateless. (such as `currentThread()`) *can be tagged and executed on FaaS*

System Design – Reducing the Performance Overhead of Fallbacks

- **Stateful connections**
- Web applications contain **stateful connections** with external services like databases. Those connections cannot be directly offloaded to FaaS platforms.
- The core idea of the **proxy-based approach** to manage stateful connections is to **share a connection to external services** between servers and FaaS functions.

System Design – Reducing the Performance Overhead of Fallbacks

- Stateful connections



System Design – Reducing the Performance Overhead of Fallbacks

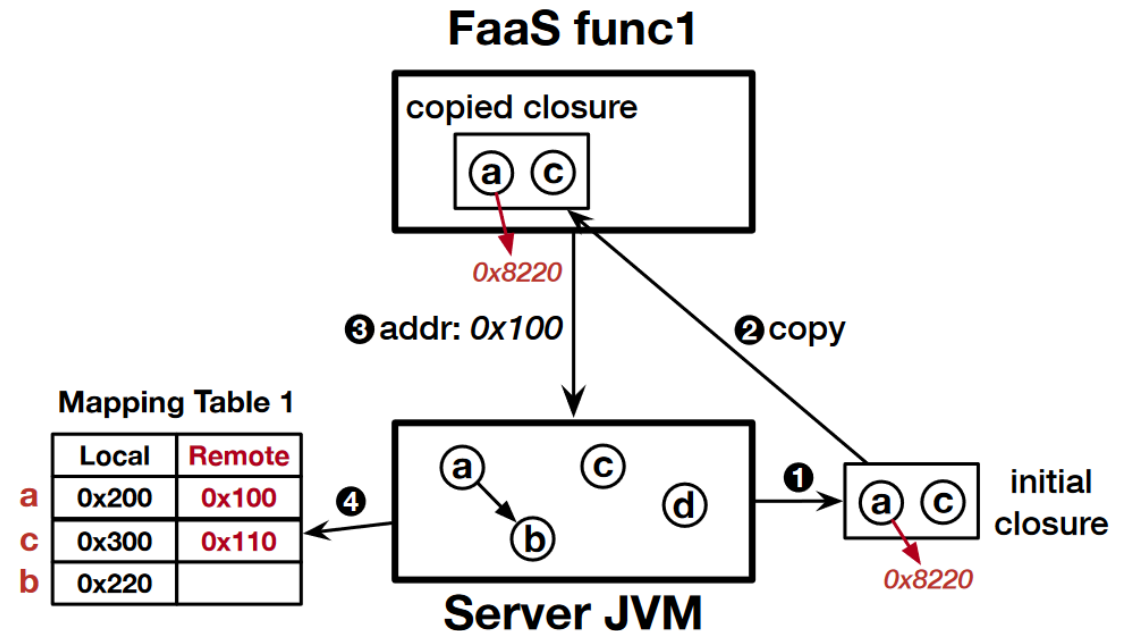
- **Warmup overhead**
- The **number of fallbacks is large** for the first execution on FaaS, and the FaaS platform needs to **establish a new runtime environment** for function execution.
- BeeHive proposes **shadow execution** to hide the warmup overhead from users, which is to process a **duplicated user request** without introducing observable state modifications.

System Design – The BeeHive Runtime System

- Laying the foundation of offloading, the BeeHive runtime is responsible for handling all communications between servers and FaaS platforms, including:
 - State management
 - Closure construction
 - Memory management

System Design – The BeeHive Runtime System

- **Distributed object sharing**
- When a function is being launched, the server constructs the **initial closure** to include objects likely to be used by the offloaded function, and mark the references of remote objects as **remote references**.

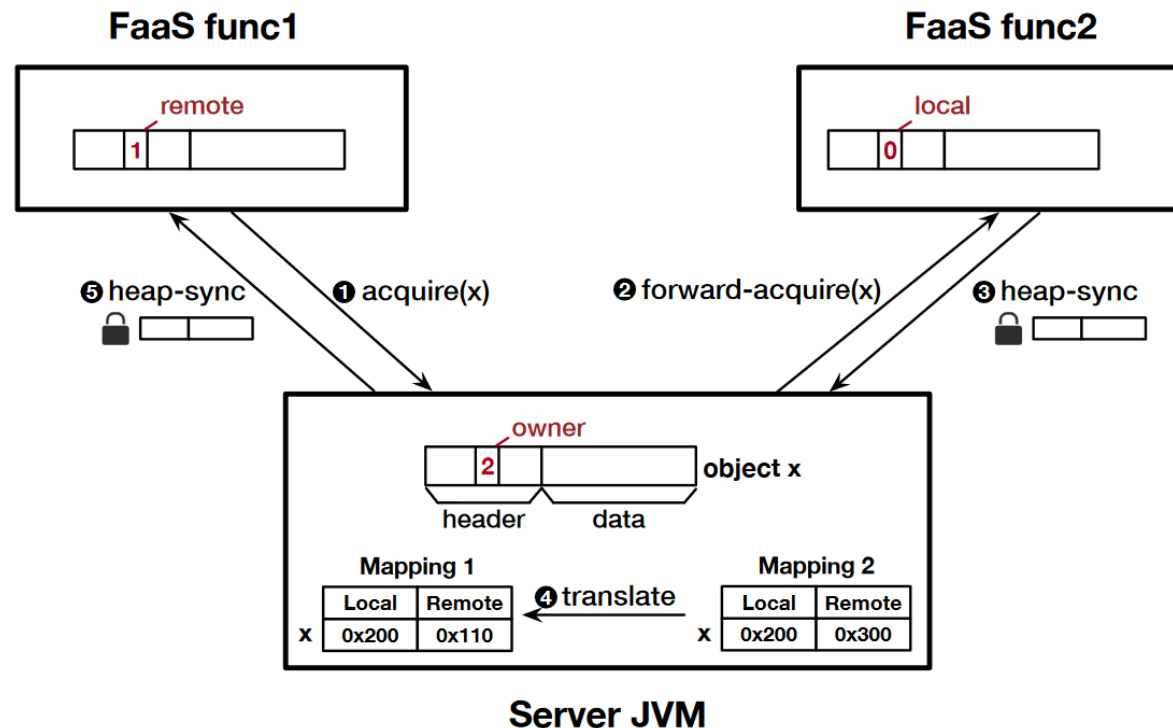


System Design – The BeeHive Runtime System

- **Shared state synchronization**
- BeeHive support **state synchronization** to ensure consistent execution for multiple FaaS endpoints.
- JMM states the **happen-before relationship** with object locks: if thread A acquires a lock previously released by thread B, then all memory operations before the lock releasing operation in thread B should be observed by thread A.

System Design – The BeeHive Runtime System

- Shared state synchronization




System Design – The BeeHive Runtime System

- **Root method selection**
- The initial closure for offloading is constructed from a **root method**.
- Method annotations can be used to distinguish **user-provided methods** from **framework ones** to avoid unsatisfying performance.

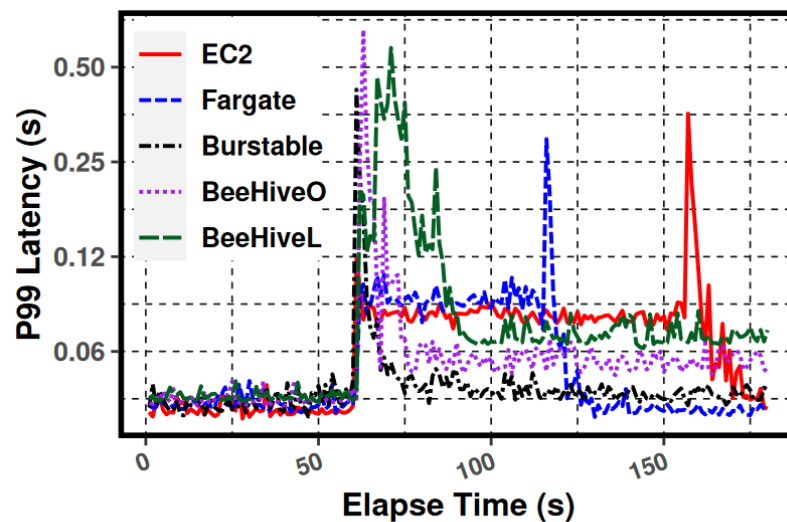
System Design – The BeeHive Runtime System

- Memory management
- Failure Recovery

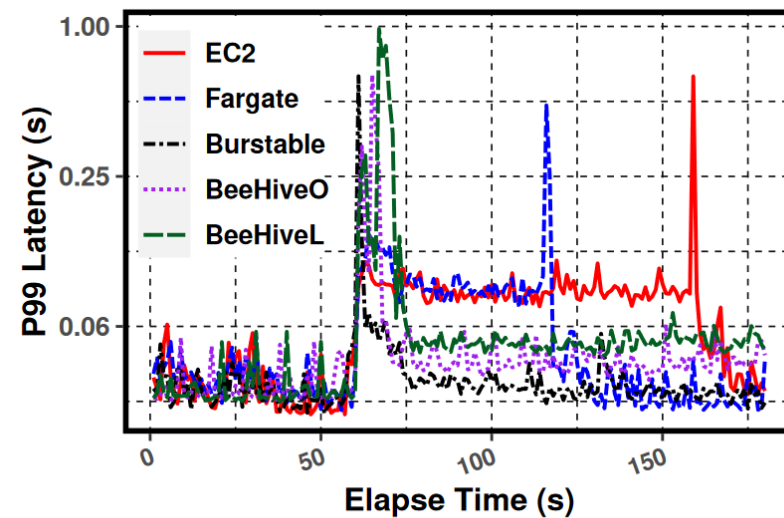
Evaluation



(a) thumbnail




(b) pybbs




(c) blog


Evaluation



(a) thumbnail



(b) pybbs



(c) blog

Evaluation

Table 3: Financial cost for scaling in Figure 7

Scaling solutions	thumbnail	pybbs	blog
EC2	0.007	0.007	0.007
Fargate	0.008	0.008	0.008
Burstable	0.005	0.005	0.005
BeeHiveO	0.010	0.017	0.013
BeeHiveL	0.012	0.010	0.008

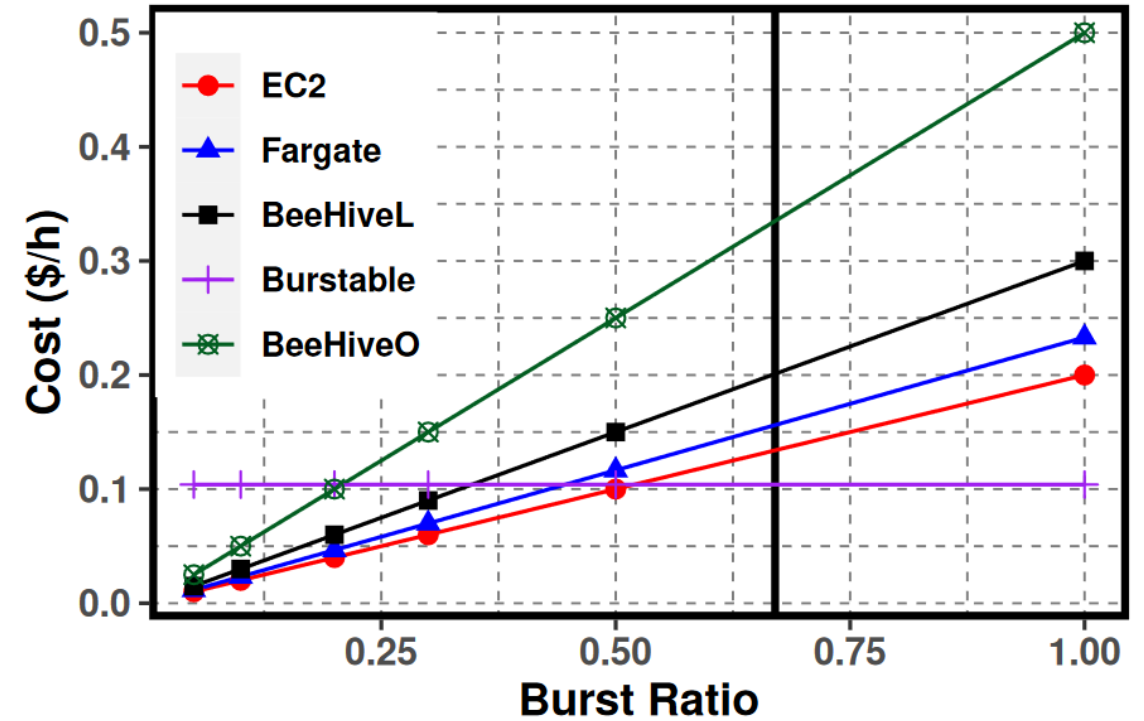


Figure 9: Cost with various burst ratios