


***BlastNet*: Exploiting Duo-Blocks for Cross-Processor Real-Time DNN Inference**

Ling, Xuan Huang, Zhihe Zhao, Nan Guan, Zhenyu Yan, and Guoliang
Xing

SenSys '22

Introduction

- **CPU-GPU heterogeneous architectures for edge platforms**




- **NVIDIA Xavier**
 - **An 8-core CPU**
 - **A Volta GPU**
- **Google Pixel 6**
 - **An 8-core CPU**
 - **A MaliG78 MP20 GPU**

Introduction


- **Cross-Processor Real-Time DNN Inference**



autonomous driving



smart roadside infrastructure



embedded computer vision

Motivation

- Platforms

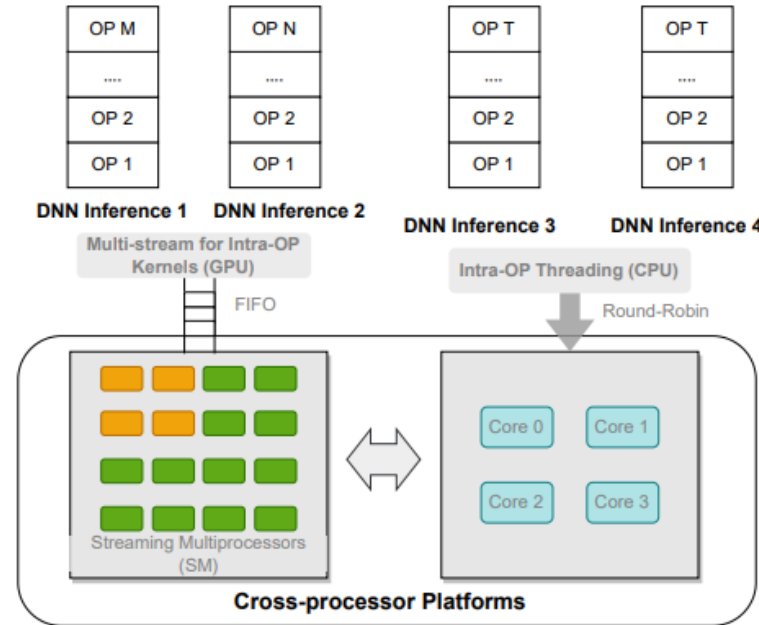


Figure 1: Concurrent DNN inference under PyTorch framework on heterogeneous CPU-GPU platforms.

Current mainstream deep learning frameworks monolithically allocate heterogeneous resources for concurrently executed DNN models..

Motivation

- **Model-level DNN Inference**





Figure 2: Worst-case execution time for concurrent DNN model inference on CPU-GPU platform under different resource allocation strategies.


Allocate DNN models to CPU and GPU in a **model-level** granularity leads to **severe resource contention**.

Motivation

- **Layer-level DNN Inference**



(a) VGG11 execution difference on CPU/GPU



(b) AlexNet execution difference on CPU/GPU

Figure 3: CPU/GPU execution ratio for each DNN layer.

Allocate DNN models to CPU and GPU in a layer-level granularity may cause **low resource utilization** and **significant layer switching overhead**.

Motivation

- **Layer-level DNN Inference**

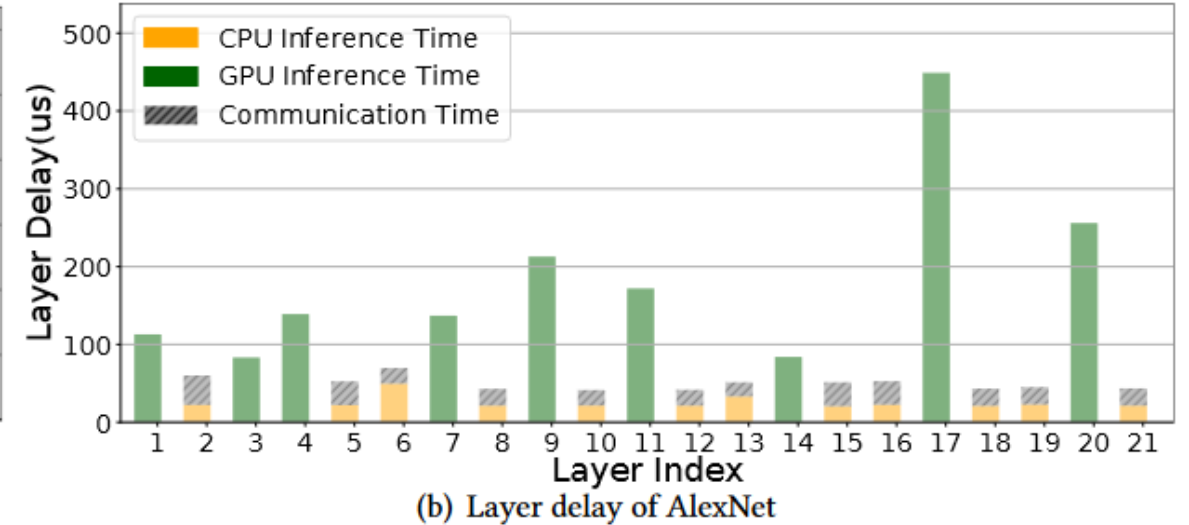




Figure 4: Layer delay by executing each layer on the processor with the shortest inference time (CPU Utilization: 9.05% for VGG11, 13.37% for AlexNet, GPU Utilization: 65.59% for VGG11, 73.39% for AlexNet, Communication Overhead: 25.35% for VGG11, 13.37% for AlexNet).

Allocate DNN models to CPU and GPU in a layer-level granularity may cause **low resource utilization** and **significant layer switching overhead**.

Design – A new abstraction of model partition

- The model-level allocation strategy often causes severe resource contention on the GPU while leaving the CPU idle.
- The layer-level allocation may lead to frequent layer switching and significant communication overhead.



Design – System Overview

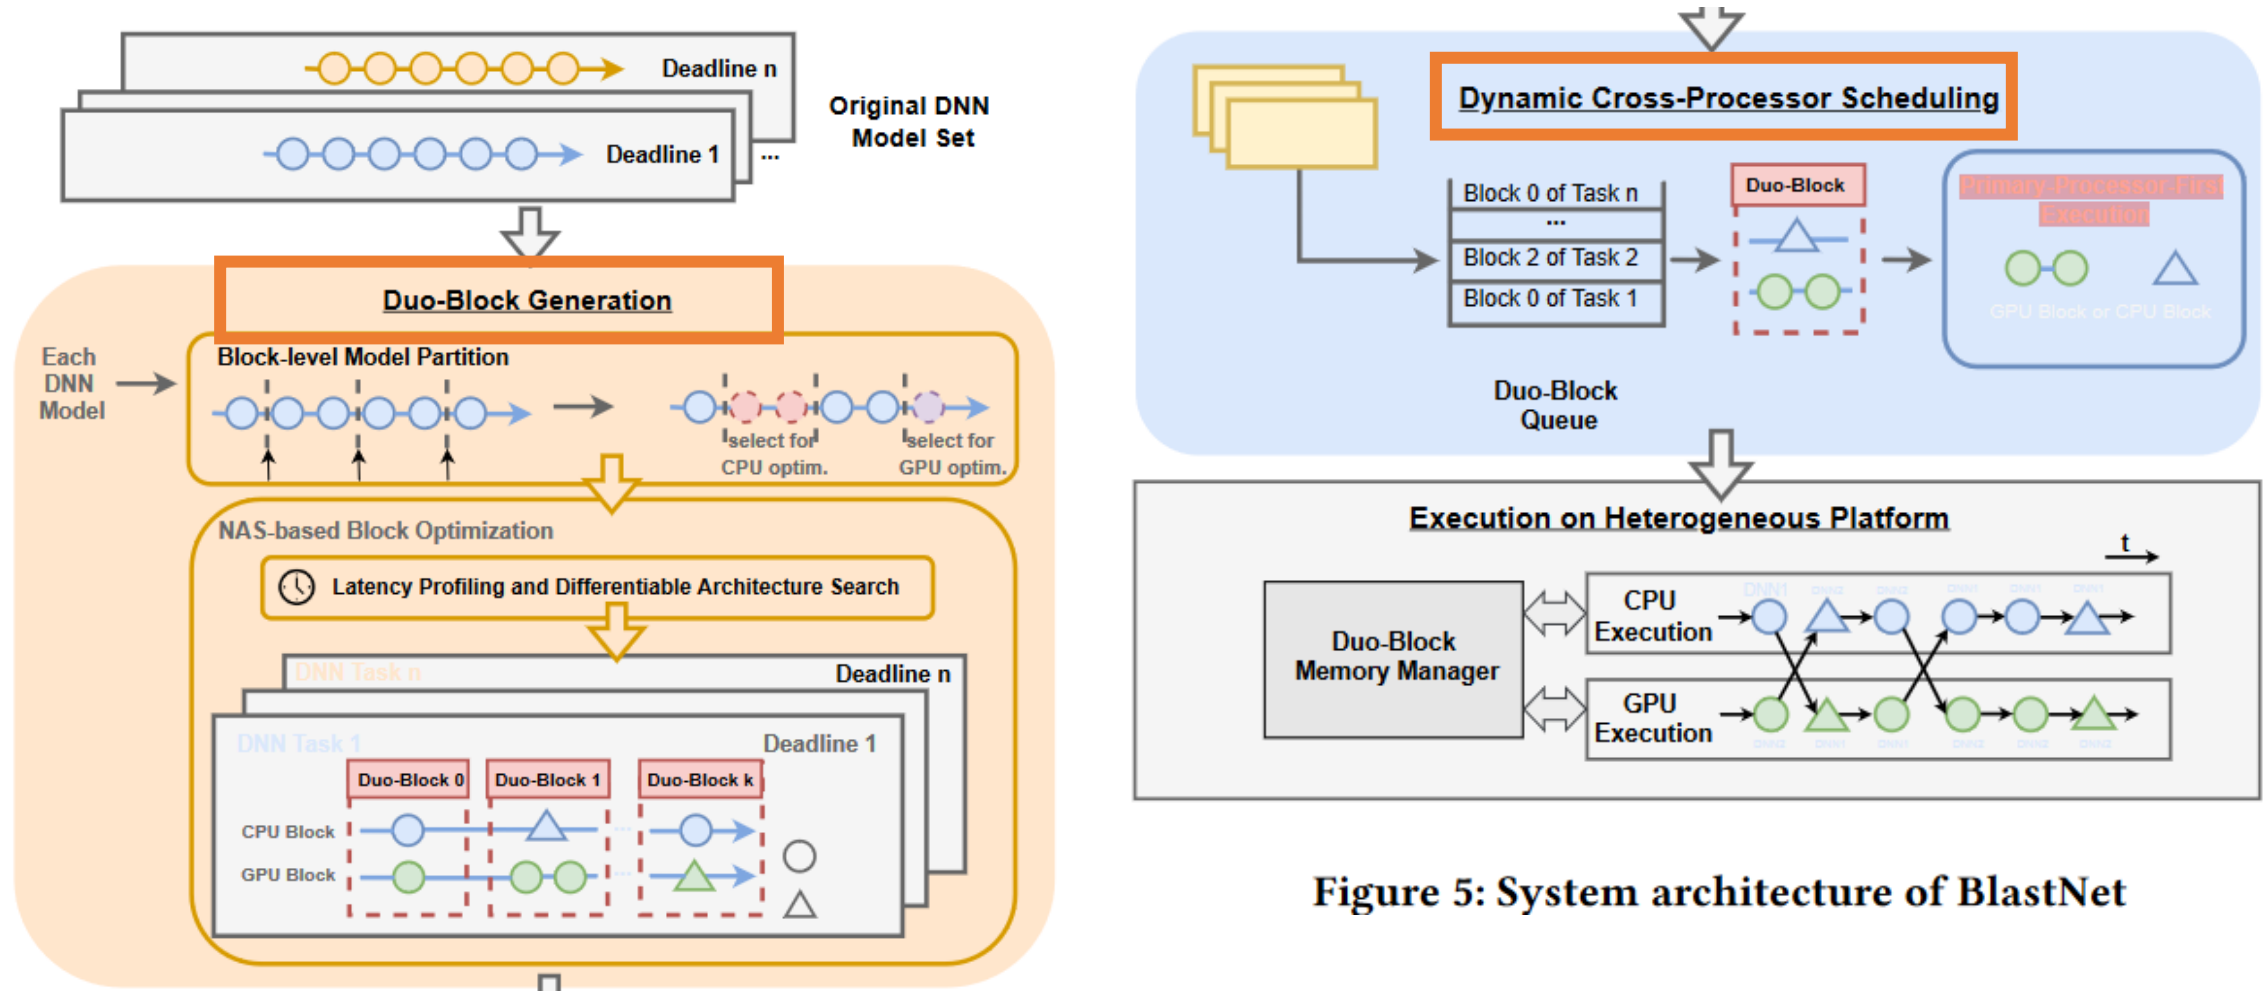
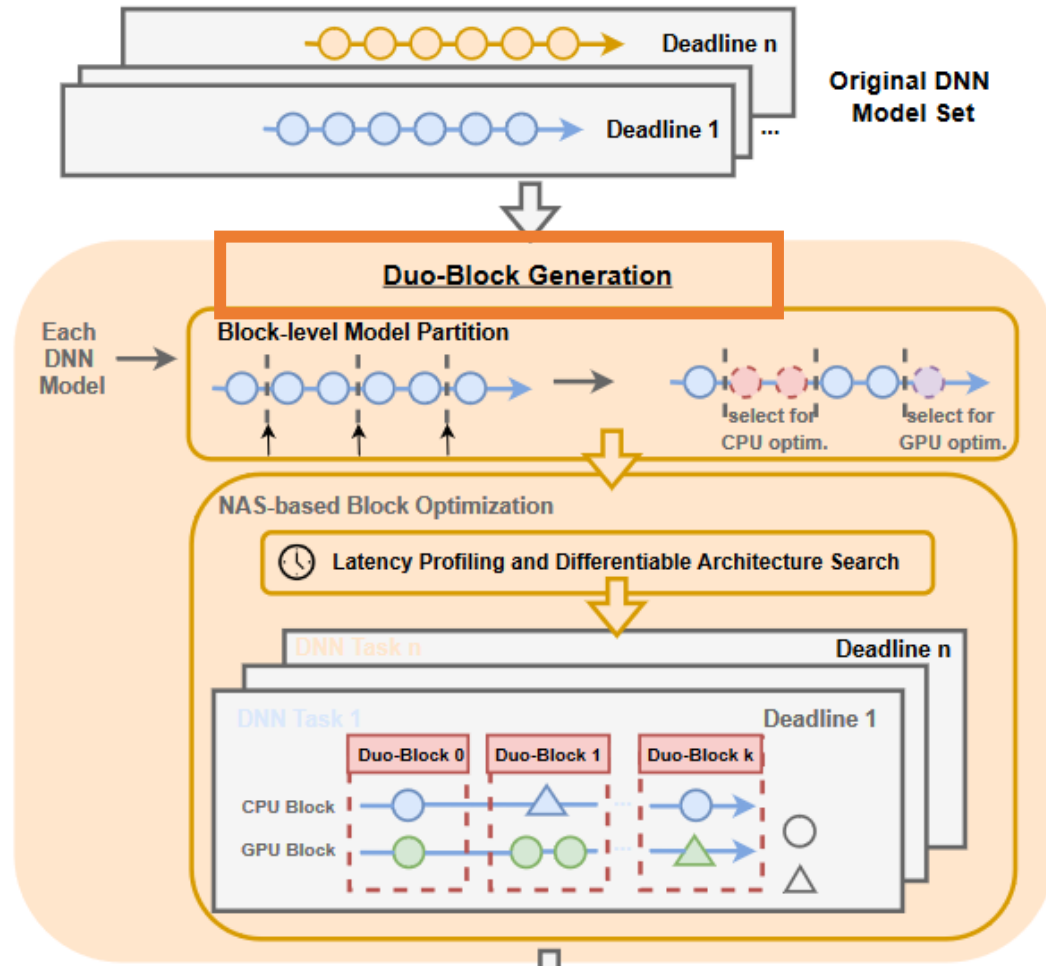


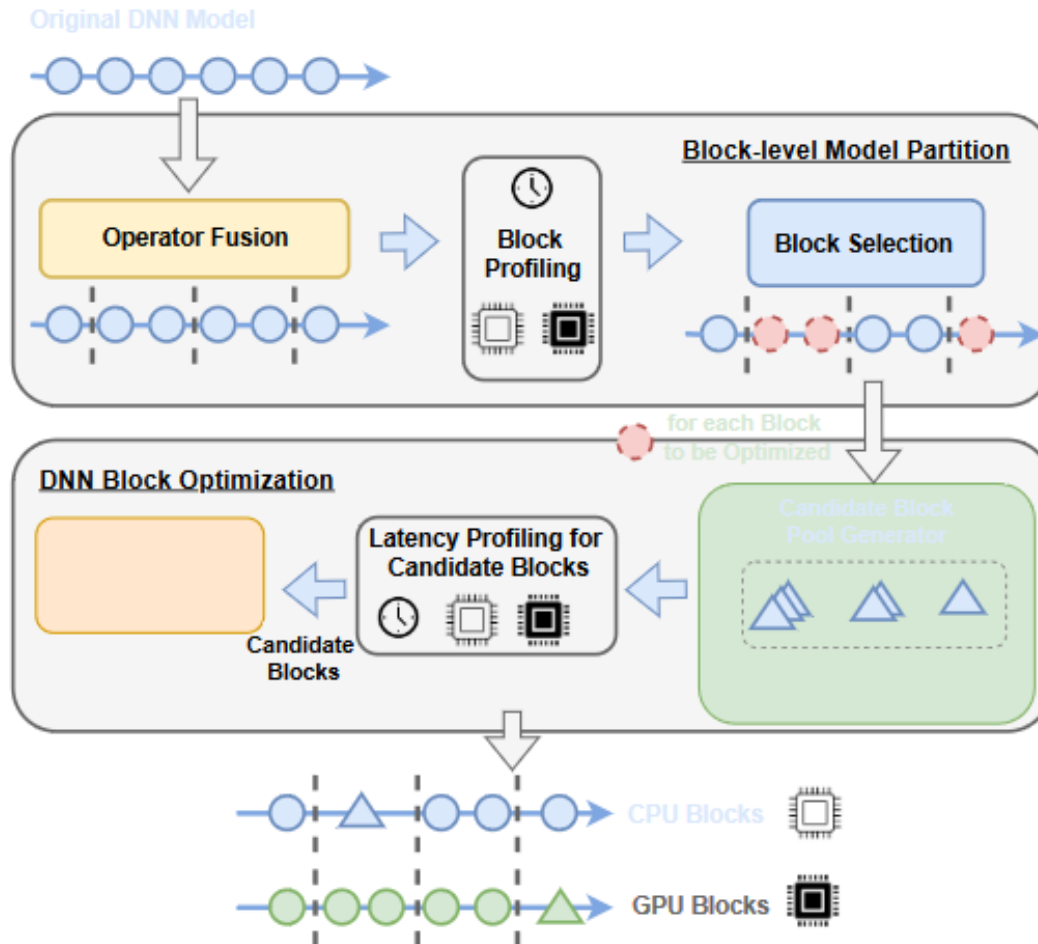
Figure 5: System architecture of BlastNet

Design – System Overview



- **Block-level Model Partition**
 - The layer-level computing
 - Communication characteristics
 - operator fusion rules

Design – duo-block generation



- **Block-level Model Partition**
 - Fuse DNN layers into blocks based on the **general operator fusion rules** and **layer characteristics**.
 - Determine the primary and the secondary processors for each block.
 - Optimize the block execution on its secondary processor.

Figure 6: Generation procedure for cross-processor block

Design – duo-block generation

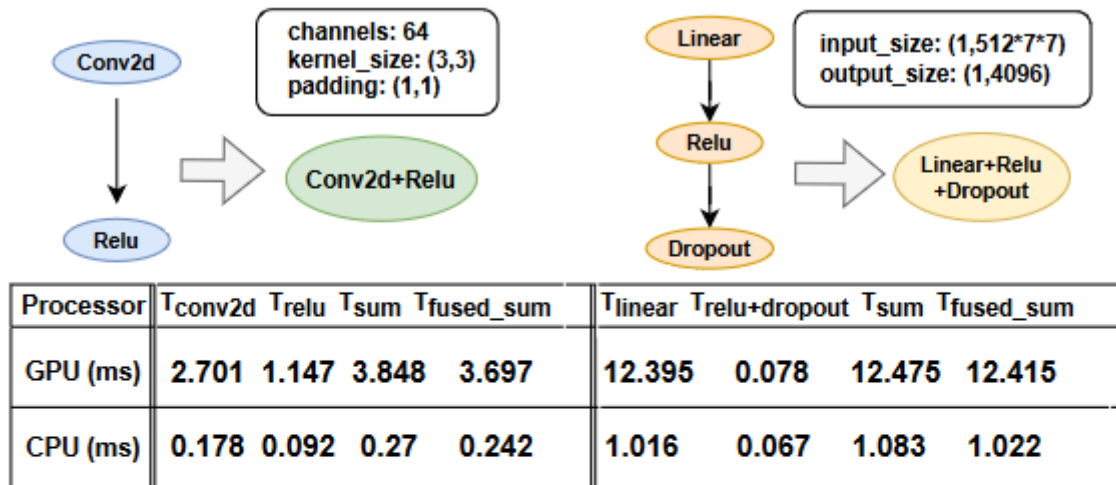
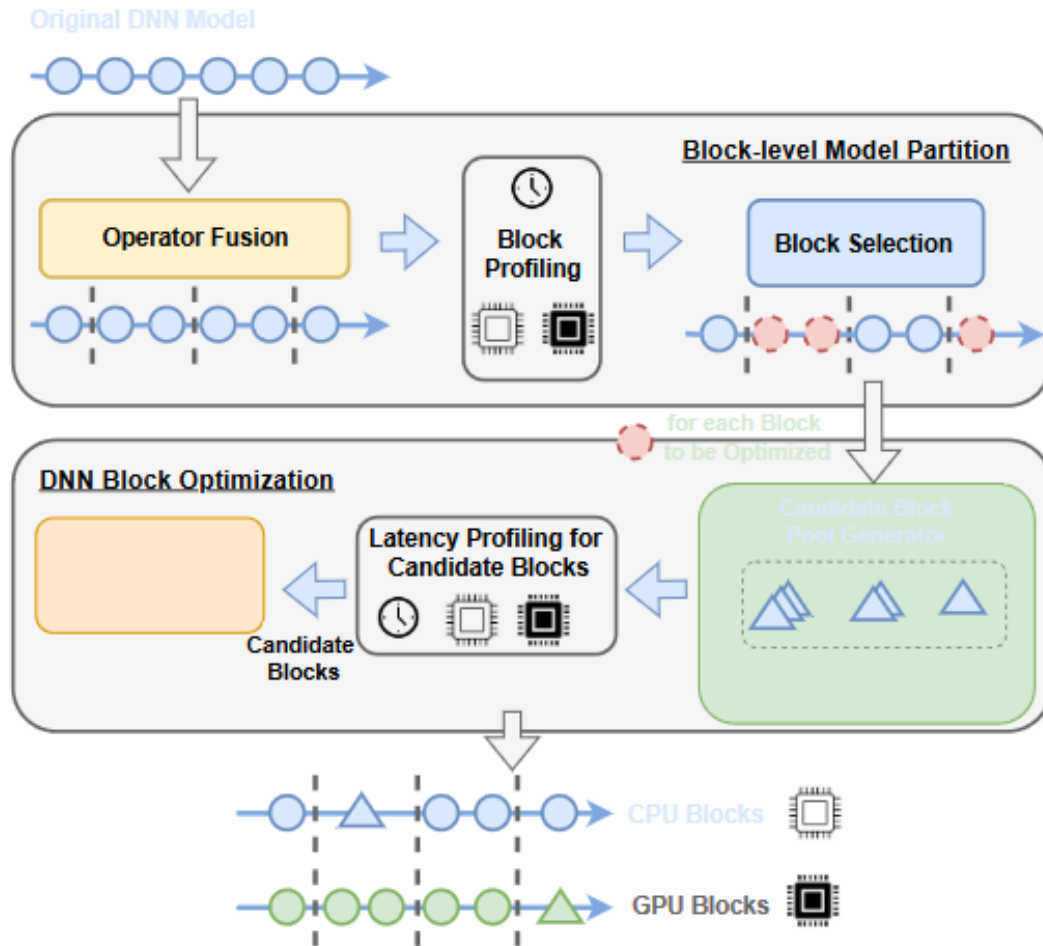


Figure 7: Operator fusion and its benefit (evaluated under torchscript on the desktop platform with NVIDIA RTX 2060 GPU). $T_{operator_name}$ denotes the execution time of the operator. T_{sum} denotes the sum of T_{conv2d} , T_{relu} or T_{linear} , $T_{dropout}$. T_{fused_sum} denotes the execution time of fused operators.

- **Block-level Model Partition**
 - Fuse DNN layers into blocks based on the **general operator fusion rules** and **layer characteristics**.
 - Determine the primary and the secondary processors for each block.
 - Optimize the block execution on its secondary processor.

Design – duo-block generation



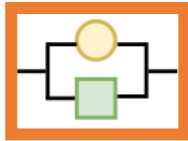
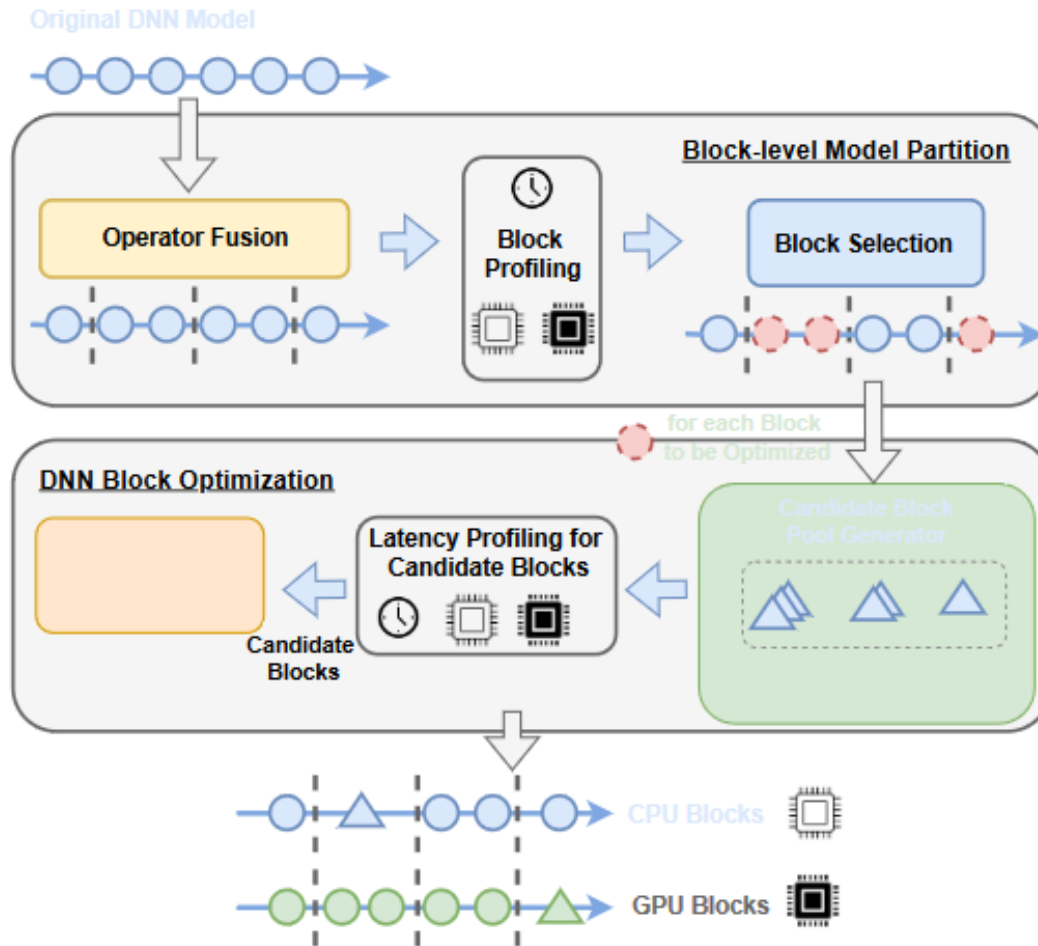
- **Block-level Model Partition**
 - Fuse DNN layers into blocks based on the general operator fusion rules and layer characteristics.
-  Determine the **primary** and the **secondary** processors for each block.
- Optimize the block execution on its secondary processor.

Figure 6: Generation procedure for cross-processor block

Design – duo-block generation



- **Block-level Model Partition**

- Fuse DNN layers into blocks based on the general operator fusion rules and layer characteristics.
- Determine the primary and the secondary processors for each block.
- **Optimize** the block execution on its secondary processor.

$$CD = \frac{T^{sec}(B_k)}{T^{pri}(B_k)}, \quad WP = \frac{T^{pri}(B_k)}{\sum_{k=0}^{k_{max}} T^{pri}(B_k)}$$

$$\hat{CD} > \varepsilon \quad WP > 1/block_num$$

Figure 6: Generation procedure for cross-processor block

Design – duo-block generation

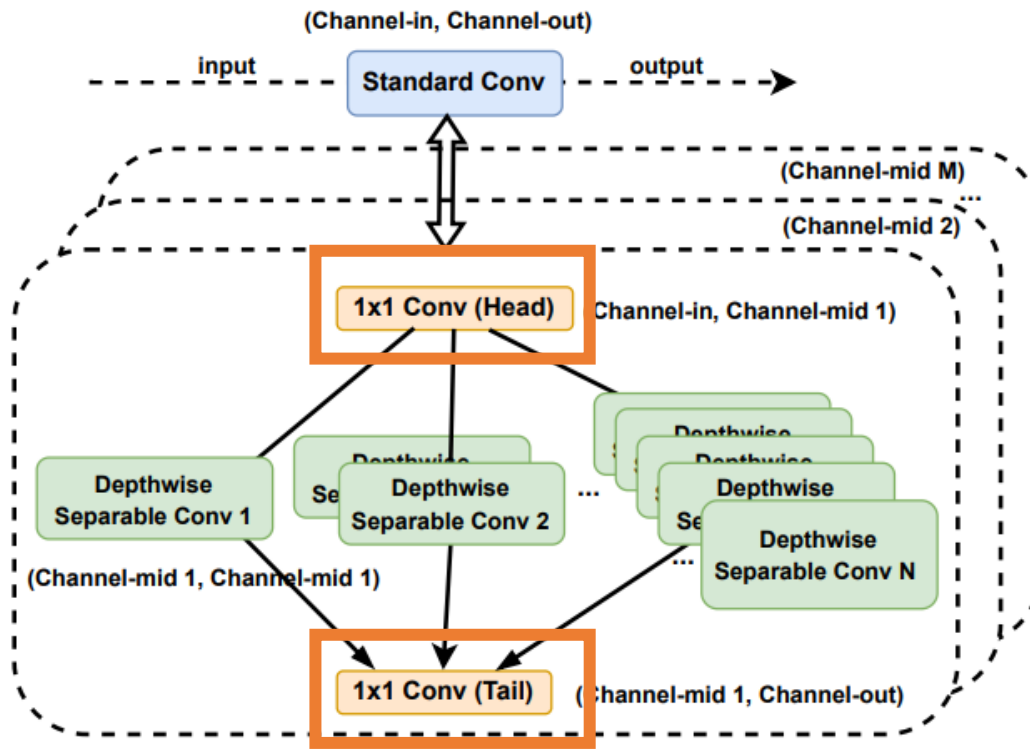


Figure 8: Example for candidate blocks of a convolutional layer on CPU

- NAS-based Block Optimization
 - Only consider **convolutional** and **fully-connected layers** since they account for the most execution time.
 - Convolutional layer optimized in CPU
 - choose the most CPU-friendly operators (i.e., **depthwise separable convolutional layer**)
 - Convolutional layer optimized in GPU
 - Fully-connected layer optimized

Design – duo-block generation

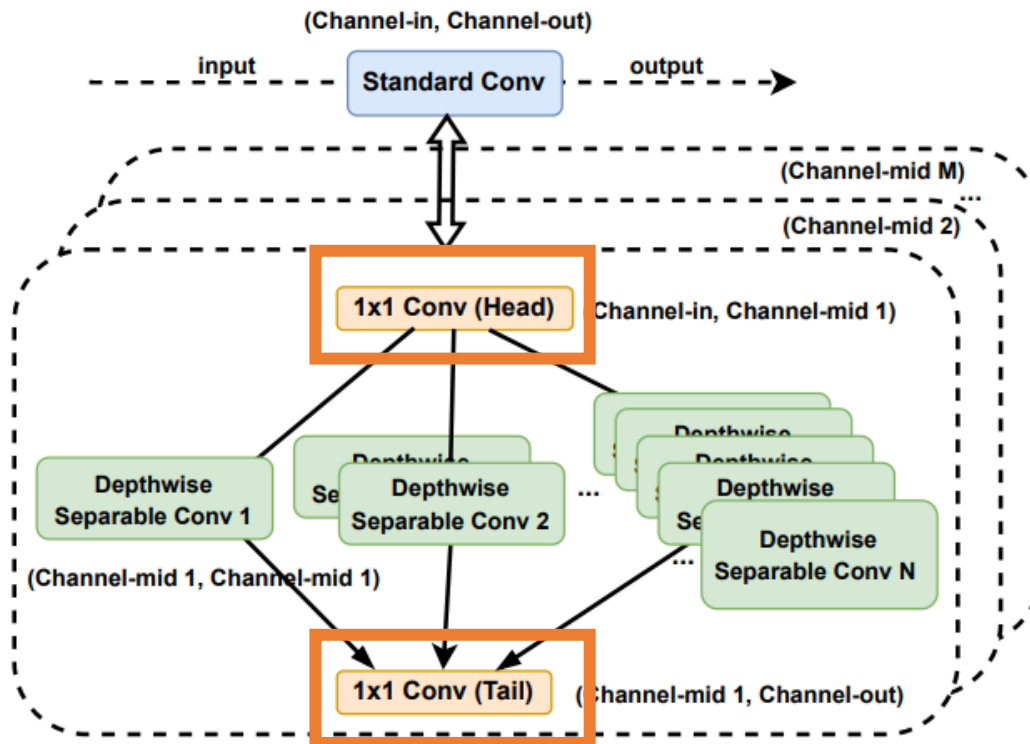


Figure 8: Example for candidate blocks of a convolutional layer on CPU

- NAS-based Block Optimization
 - Only consider **convolutional** and **fully-connected layers** since they account for the most execution time.
 - Convolutional layer optimized in CPU
 - Convolutional layer optimized in GPU
 - adopt **denser convolutional layers**
 - Fully-connected layer optimized

Design – duo-block generation




Figure 8: Example for candidate blocks of a convolutional layer on CPU

- NAS-based Block Optimization
 - Only consider **convolutional** and **fully-connected layers** since they account for the most execution time.
 - Convolutional layer optimized in CPU
 - Convolutional layer optimized in GPU
 - Fully-connected layer optimized
 - choose the fully-connected layer with **a smaller/larger** channel size

Design – duo-block generation

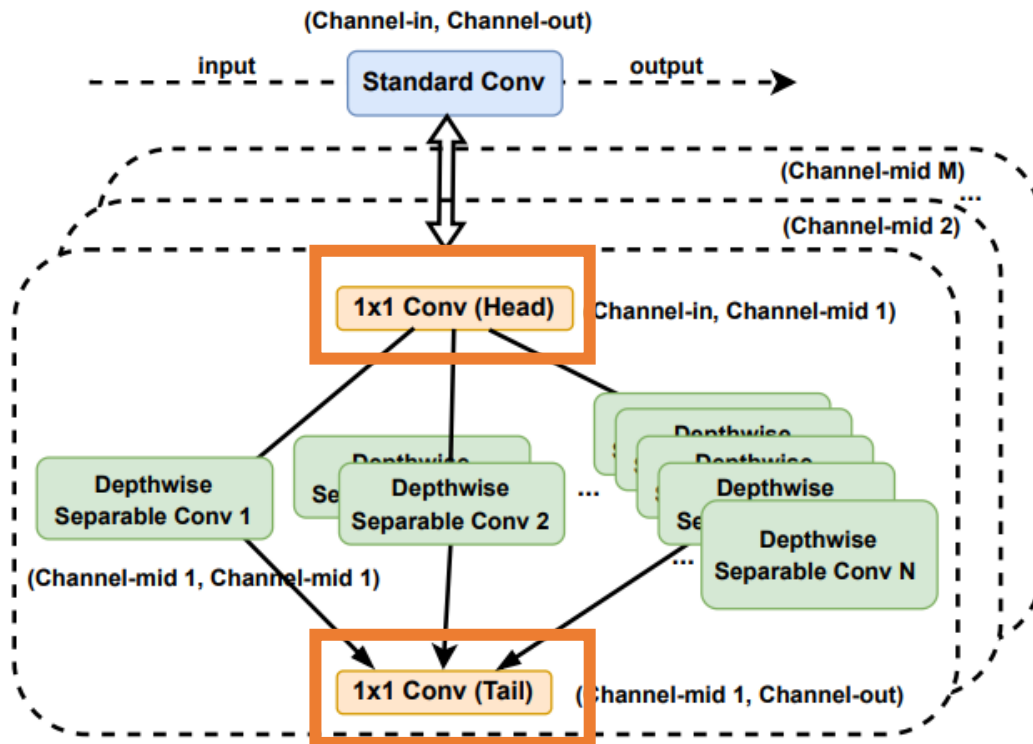
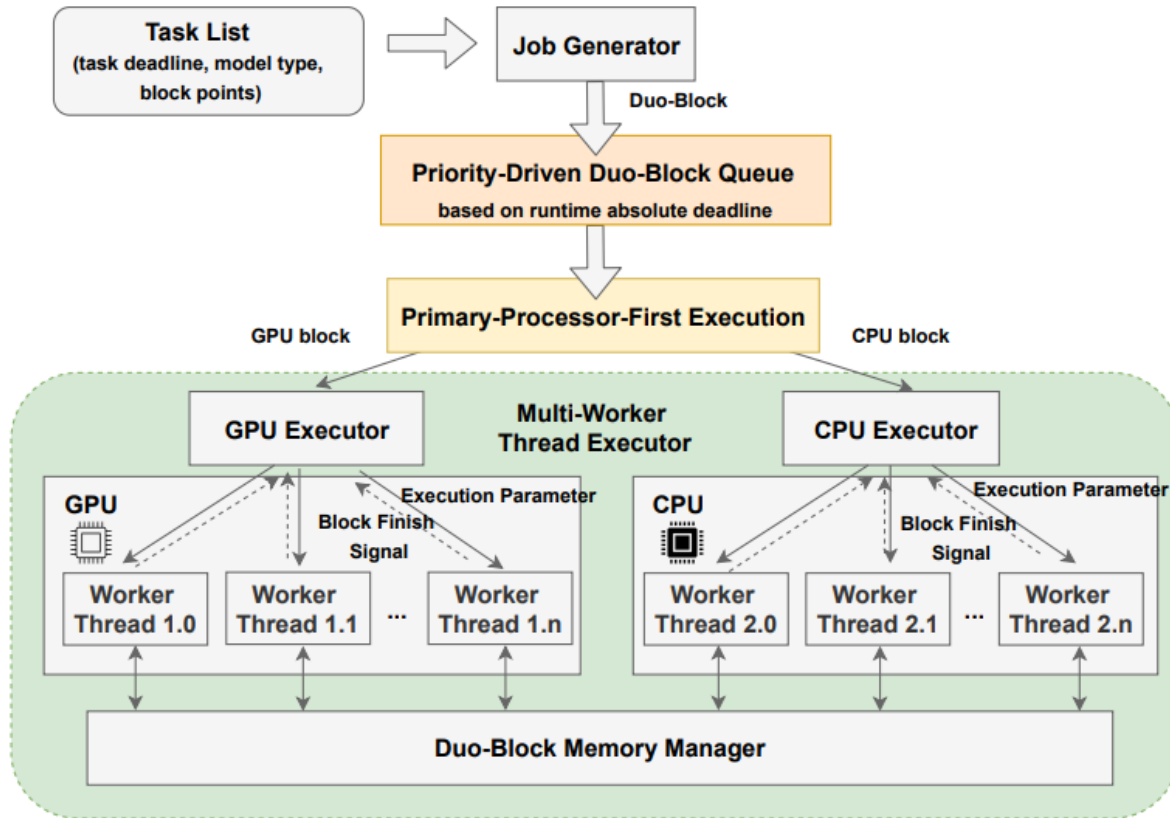


Figure 8: Example for candidate blocks of a convolutional layer on CPU

- NAS-based Block Optimization
 - Each candidate block has a **head module** and a **tail module**.
 - Search for the optimal block that minimizes the **accuracy loss**.
 - The search algorithm is based on the **Differentiable Architecture Search (DARTS)** algorithm

$$\begin{aligned} \forall k \quad & \min \mathcal{L}(B_k^{new}, W^*) \\ \text{s.t.} \quad & T^{sec}(B_k^{old}) > T^{sec}(B_k^{new}) \end{aligned}$$

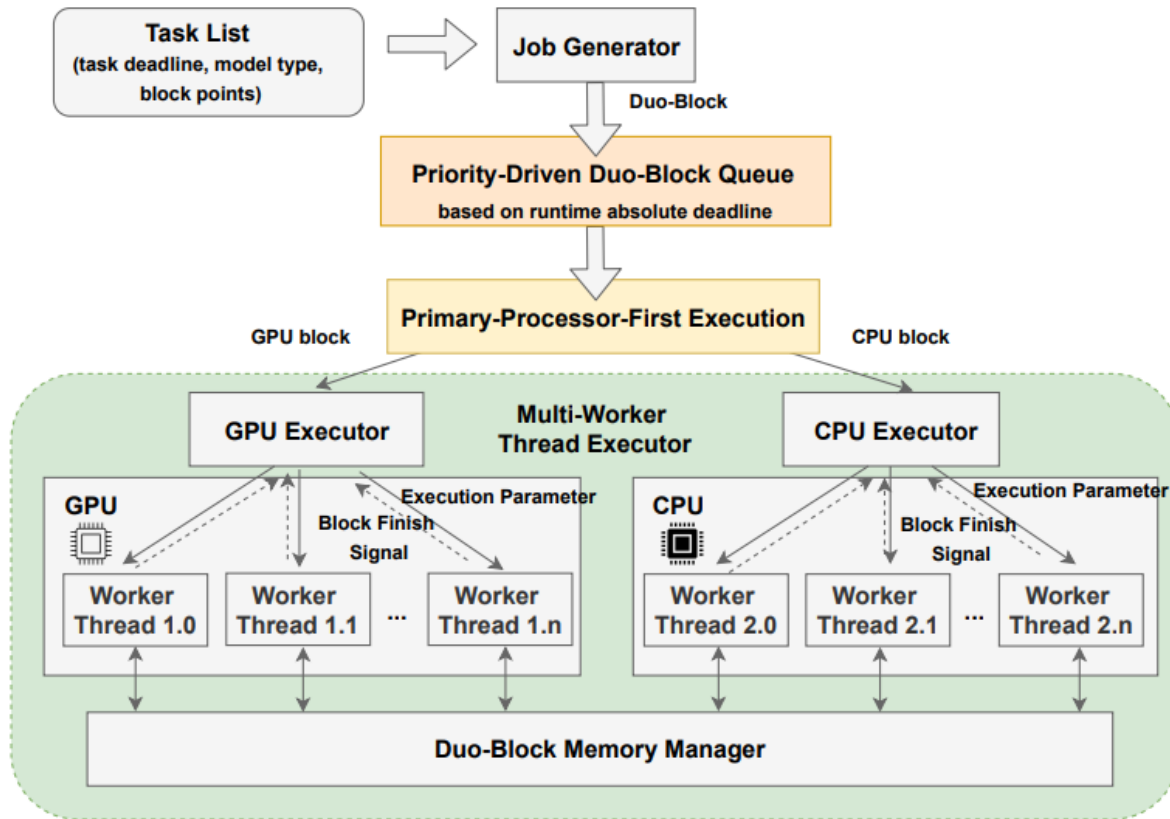
Design – dynamic cross processor scheduling



- The scheduler prioritizes each duo-block based on its **task urgency**.
- Primary processor-first execution mechanism decides the execution processor for each duo-block based on the status of the processors.

Figure 9: Procedure for cross-processor scheduling, worker thread 1.1 represents the worker thread for DNN model 1 on the GPU processor.

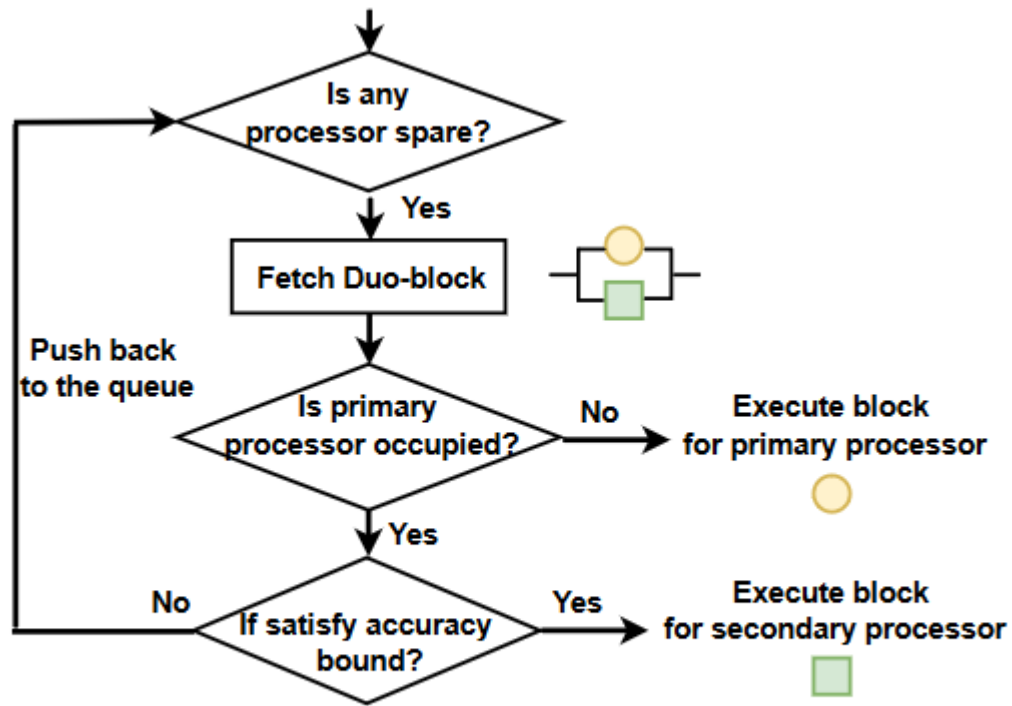
Design – dynamic cross processor scheduling



- The scheduler prioritizes each duo-block based on its task urgency.
- **Primary processor-first execution mechanism decides the execution processor for each duo-block based on the status of the processors.**

Figure 9: Procedure for cross-processor scheduling, worker thread 1.1 represents the worker thread for DNN model 1 on the GPU processor.

Design – dynamic cross processor scheduling



- The scheduler prioritizes each duo-block based on its task urgency.
- **Primary processor-first execution mechanism decides the execution processor for each duo-block based on the status of the processors.**

Figure 10: Primary-processor-first execution mechanism

Implementation

- **Platforms**

Table 1: Platforms used in evaluation experiments.

Platform	GPU	CPU	Memory	Storage
NVIDIA AGX Xavier	512-core Volta	8-core ARMv8.2	16GB	32GB
NVIDIA Jetson TX2	256-core Pascal	2-core ARM Denver + 4-core ARM A57	8GB	32GB
Desktop	RTX2080	8-core Intel i9-9900K	32GB	5TB

Implementation

- Implementation

DNN Inference Task Type	Dataset	DNN Model
Image Classification	CIFAR10 [24]	MobileNet[17], VGG11, AlexNet
Sign Recognition	GTSRB [48]	ResNet18
Object Detection	COCO [32]	YOLO [44]

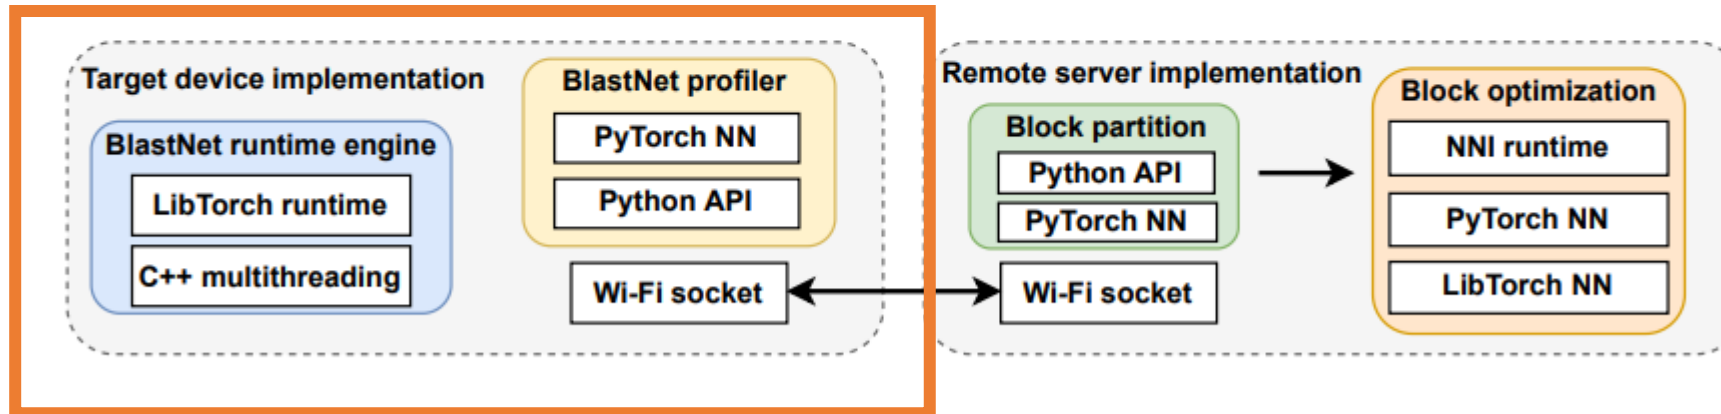


Figure 11: BlastNet Software Implementation

Implementation

- Implementation

DNN Inference Task Type	Dataset	DNN Model
Image Classification	CIFAR10 [24]	MobileNet[17], VGG11, AlexNet
Sign Recognition	GTSRB [48]	ResNet18
Object Detection	COCO [32]	YOLO [44]

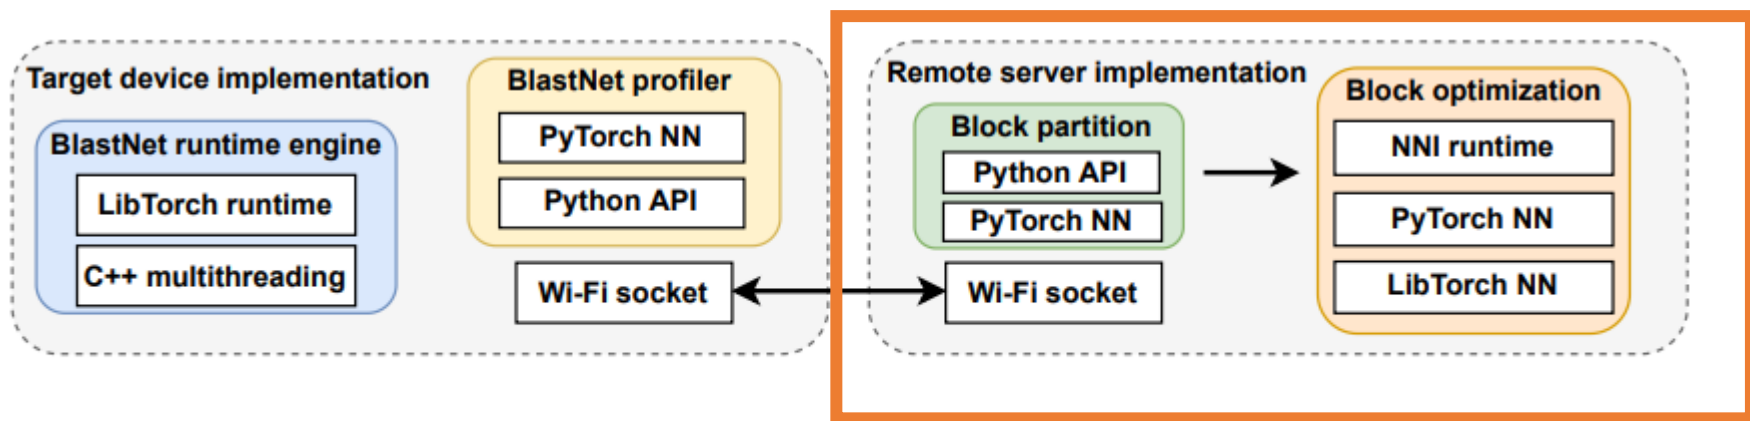


Figure 11: BlastNet Software Implementation

End-to-end System Evaluation

- **Autonomous driving testbed**



(a) F1/10 Vehicle




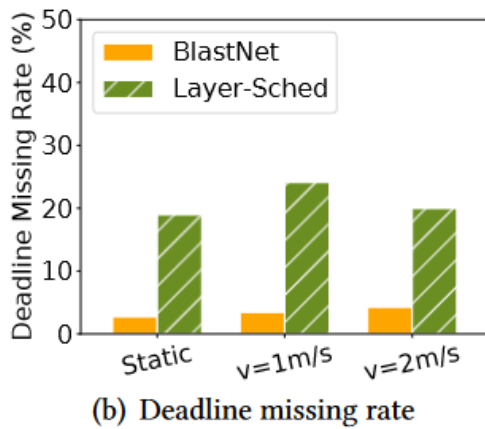
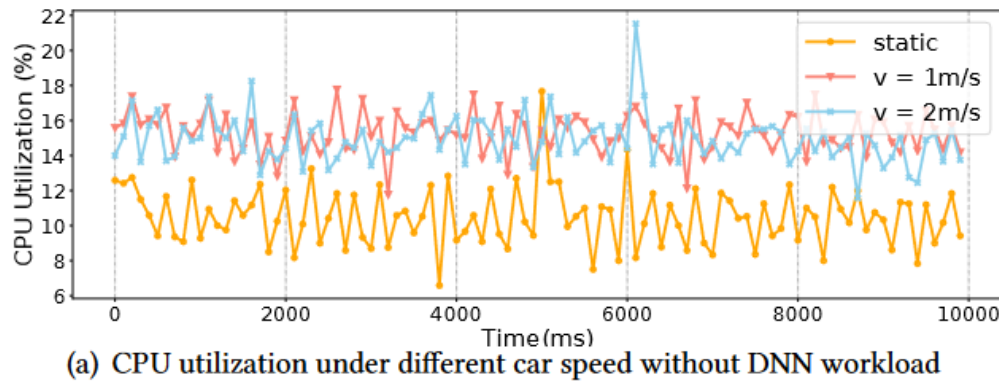
(b) Testbed setup

Figure 12: F1/10 autonomous driving testbed.

Four real-time DL tasks for **traffic sign recognition** and **a lane detection task** running on this platform.

End-to-end System Evaluation

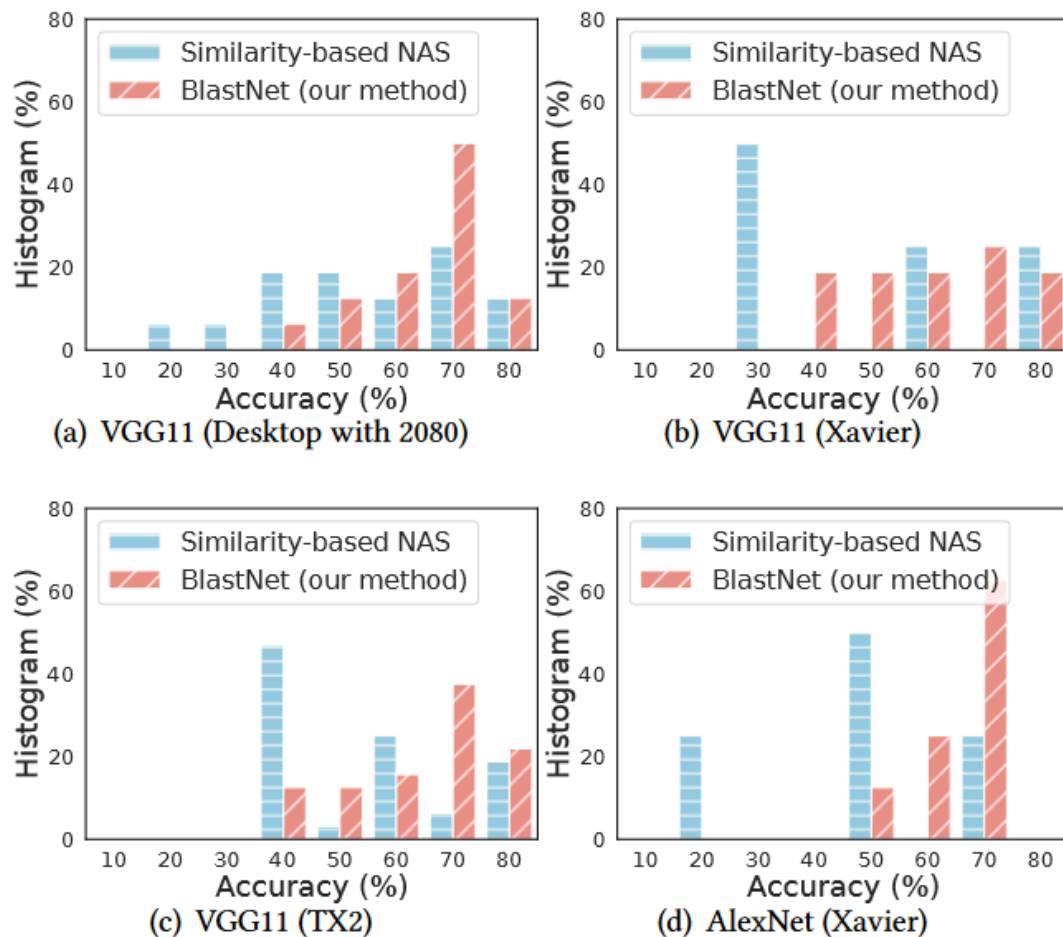
- **Autonomous driving testbed**



- **Baseline:** *Layer _Sched* schedules the DNN model at the **layer level**
- **Running the F1/10 autonomous vehicle at different speeds, resulting in different levels of resource utilization.**

Figure 13: Performance of BlastNet under various driving settings.

Performance of Cross-processor Duo-block Generation

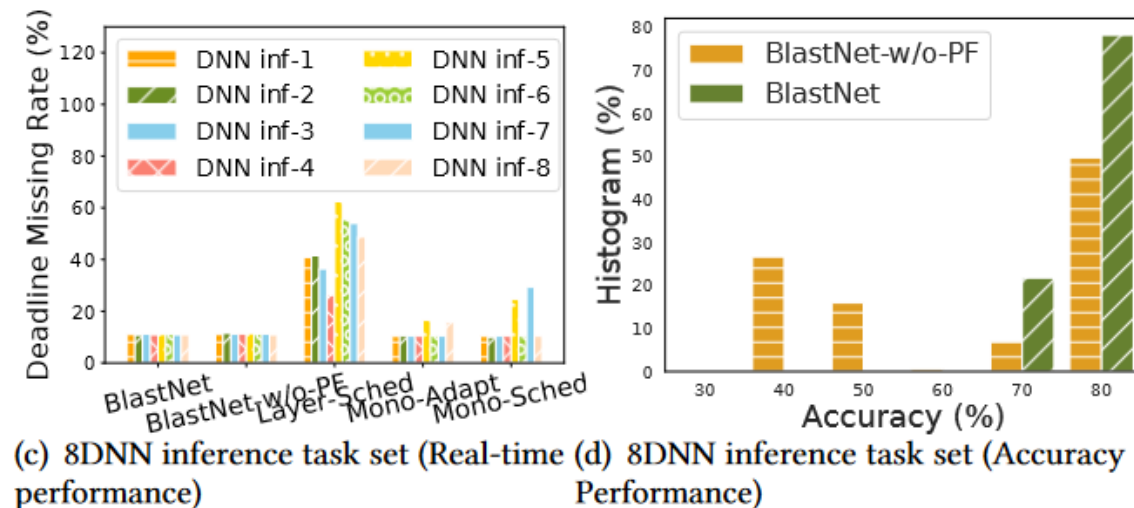
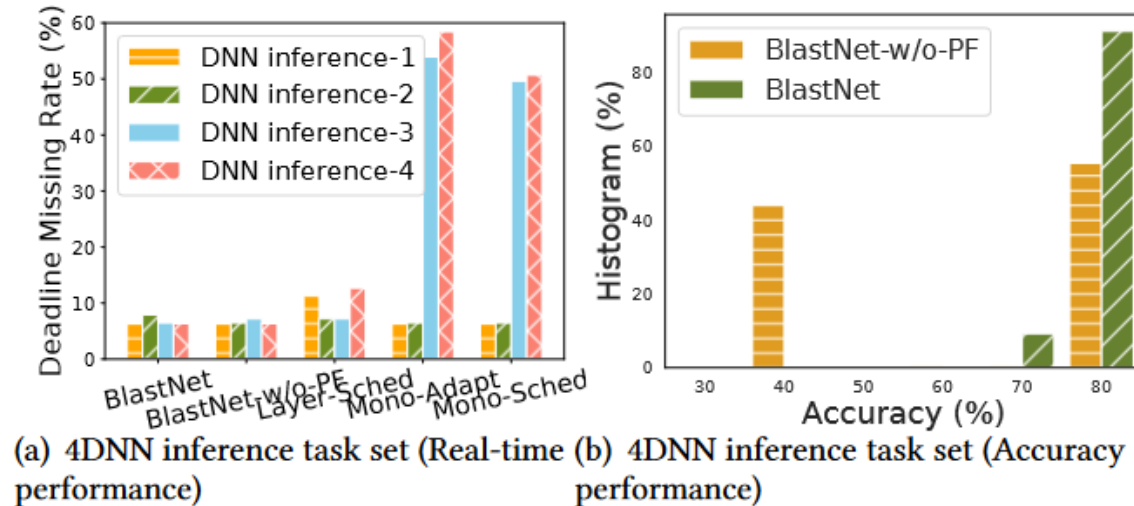


- Baseline: *Similarity-based NAS*

Figure 14: Model accuracy with all possible inference paths with the optimized blocks.

Overall Performance

• Impact of different DNN workloads



- ***Mono_Sched*** : Monolithically allocate the DNN models to heterogeneous resources.
- ***Layer_Sched*** : Schedule the DNN model at the layer level
- ***Mono_Adapt*** : Differs from BlastNet only in DNN scheduling by adopts monolithic scheduling to execute the models.
- **BlastNet w/o-PF**: Differ from BlastNet only in that it has no primary-first execution mechanism

Overall Performance

- **Impact of different DNN workloads**

	Minimum	Average	1/4 Value	Maximum
BlastNet (4task)	73.05%	80.08%	80.79%	80.79%
BlastNet-w/o-PF (4task)	46.65%	65.69%	46.65%	80.79%
BlastNet (8task)	73.05%	79.16%	80.79%	80.79%
BlastNet-w/o-PF (8task)	46.65%	67.02%	46.65%	80.79%

(e) Statistics for accuracy

Figure 15: Real-time/Accuracy performance of BlastNet under different DNN workloads.

Overall Performance

- Impact of background load


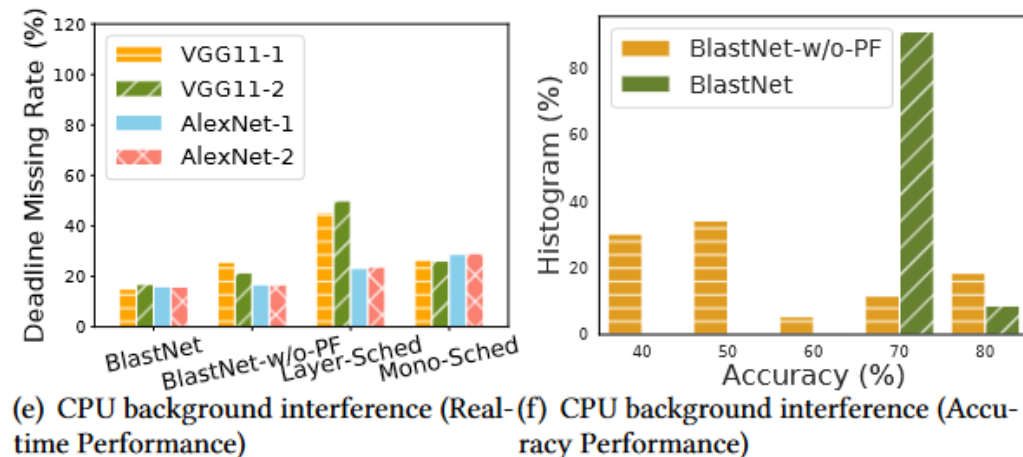
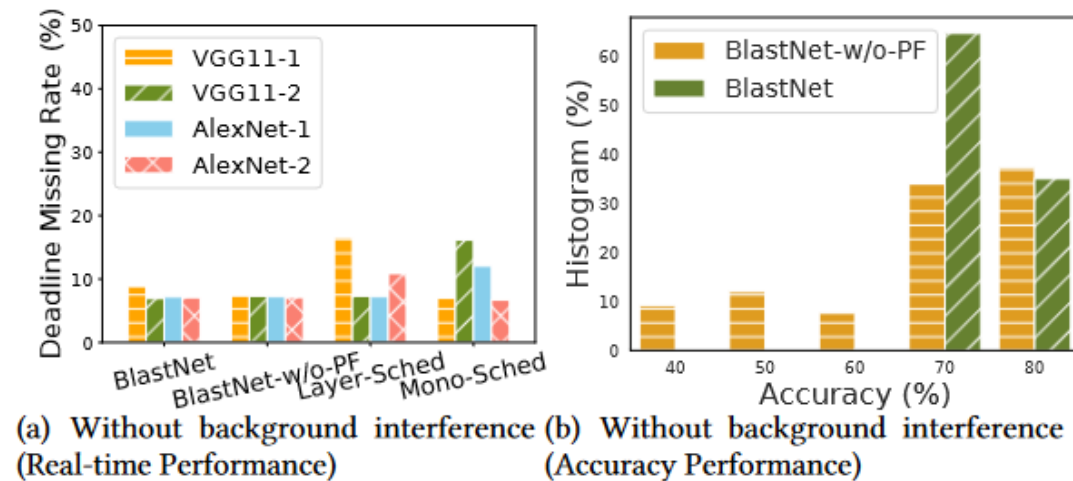




Figure 16: Real-time/Accuracy performance of BlastNet under interference.

Overall Performance

- Different edge platforms



(a) AGX Xavier (Real-time performance) (b) AGX Xavier (Accuracy performance)



(c) Jetson TX2 (Real-time performance) (d) Jetson TX2 (Accuracy Performance)

Figure 17: Performance comparison of BlastNet under different edge platforms.

Overall Performance

- CPU/GPU utilization**

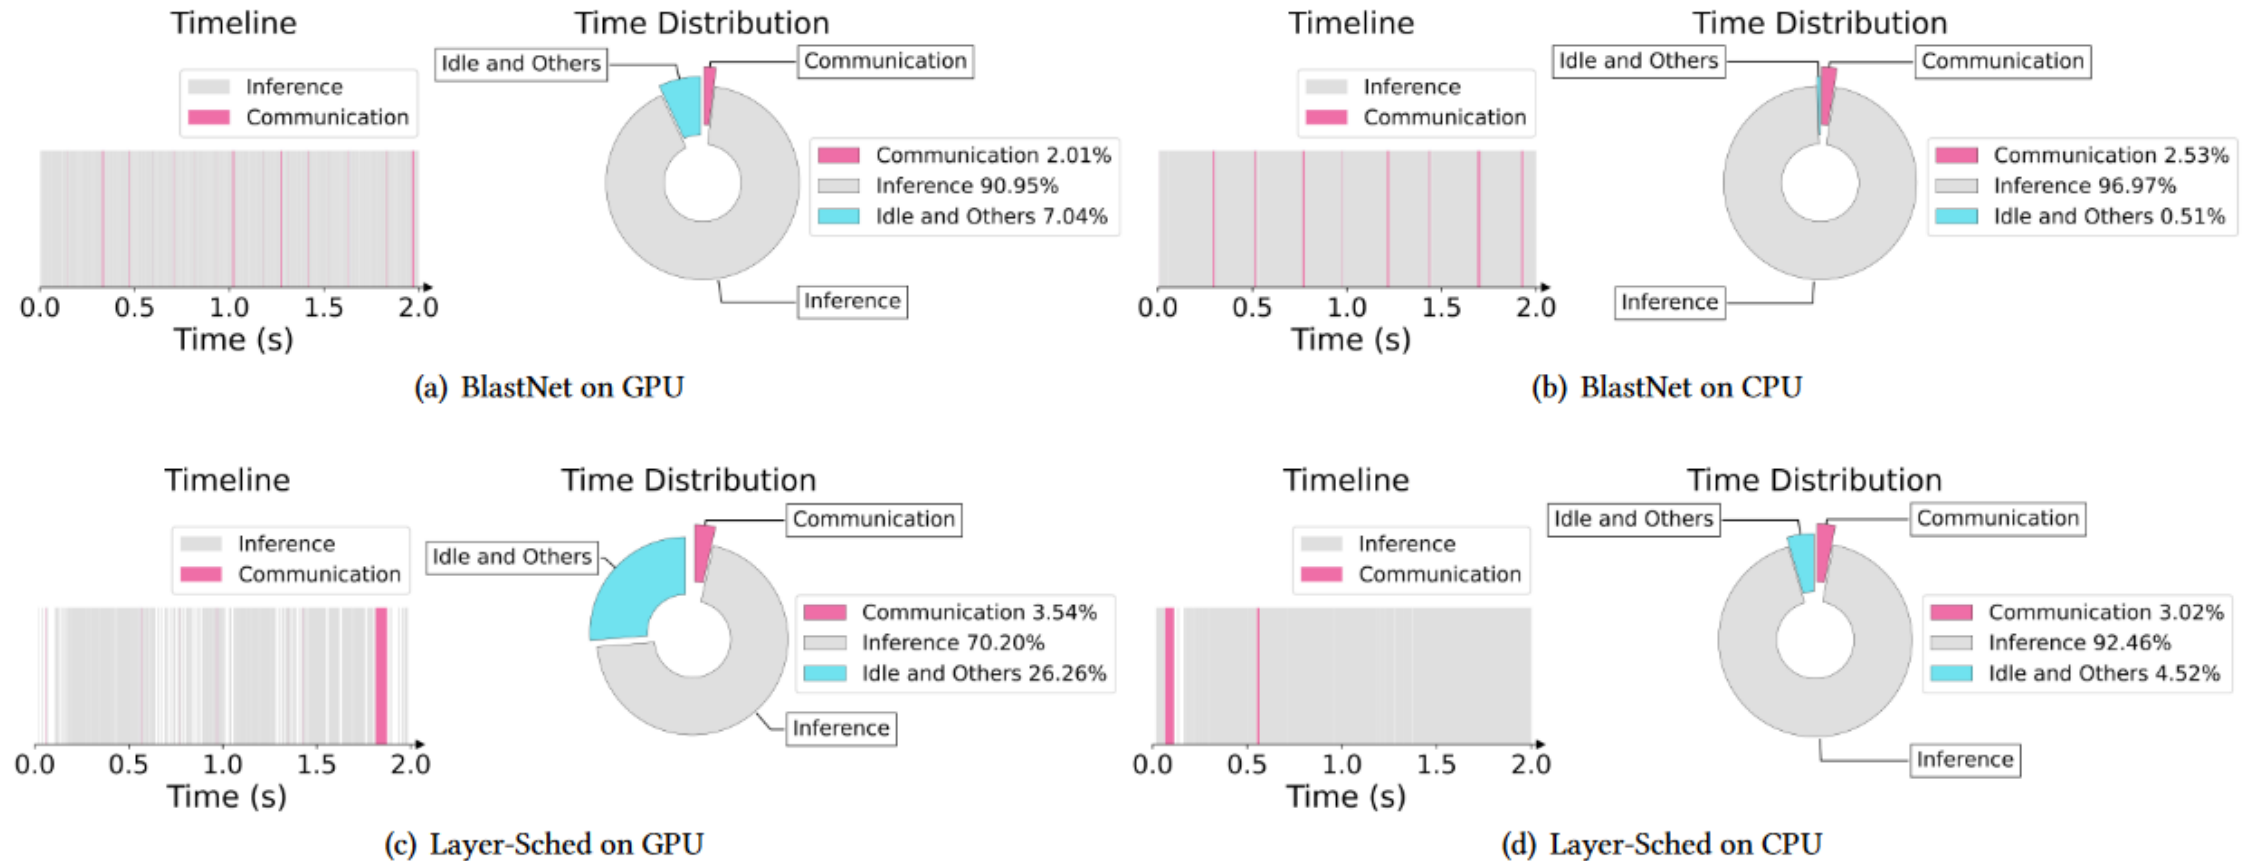


Figure 18: Execution timeline for inference, communication, and idle.

System Overhead

Table 3: CPU Overhead of Block-level DNN Scheduling

Task Set Size	Desktop	Xavier	TX2
2 DNN inference tasks	1.21%	1.61%	2.86%
4 DNN inference tasks	1.52%	1.38%	3.97%
6 DNN inference tasks	2.08 %	1.71%	3.20%

- caused by the dynamic cross-processor DNN scheduler

Advantage

- Propose a new abstraction of model partition : duo-block.
- **Efficient utilize** the resource of GPU and CPU.

Disadvantage

- The online scheduling strategy is **overly simplistic**.